

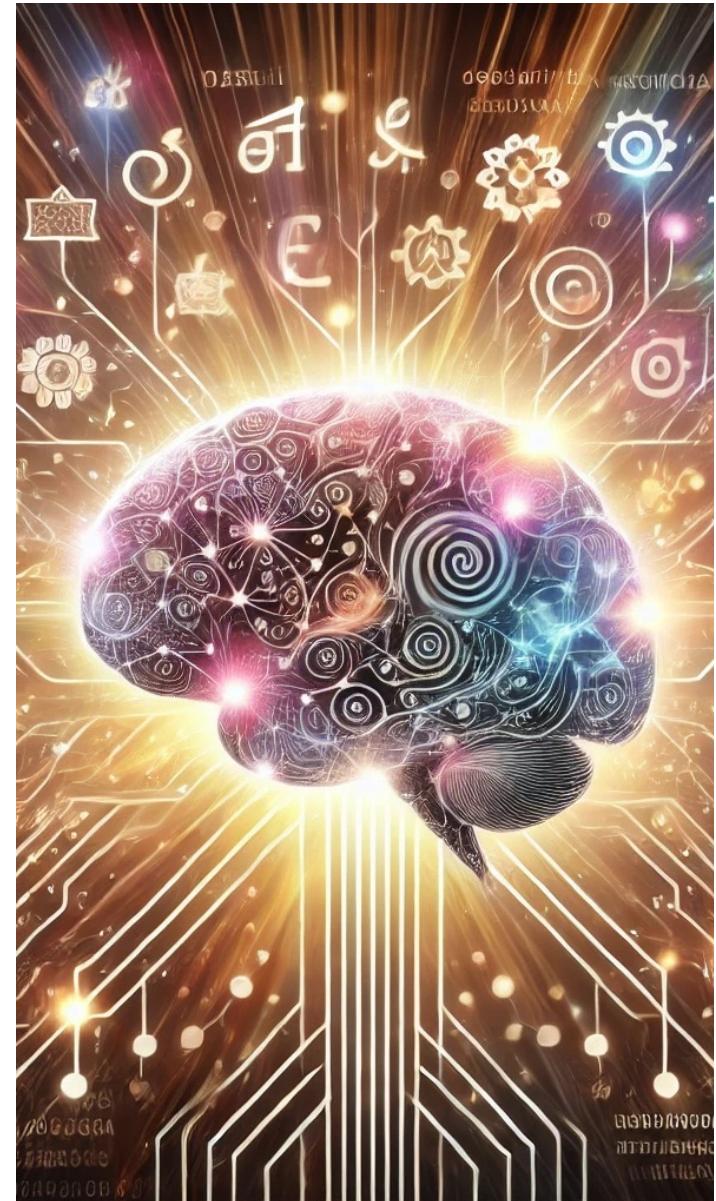
Chapter 10:

Time Series Analysis

Prachya Boonkwan (Arm)

School of ICT, SIIT,
Thammasat University

prachya@siit.tu.ac.th, kaamanita@gmail.com





<https://tinyurl.com/5b6ch5px>

Who? Me?

- Nickname: **Arm** (P'/N' Arm, etc.)
- Born: Aug 1981
- Work
 - Researcher at NECTEC 2005-2024
 - Lecturer at SIIT, Thammasat University 2025-now
- Education
 - B.Eng & M.Eng, CPE Kasetsart University
 - Obtained Ministry of Science Scholarship in early 2008
 - Did a PhD in Informatics (AI & Computational Linguistics) at University of Edinburgh, UK from 2008 to 2013 (4.5 years)



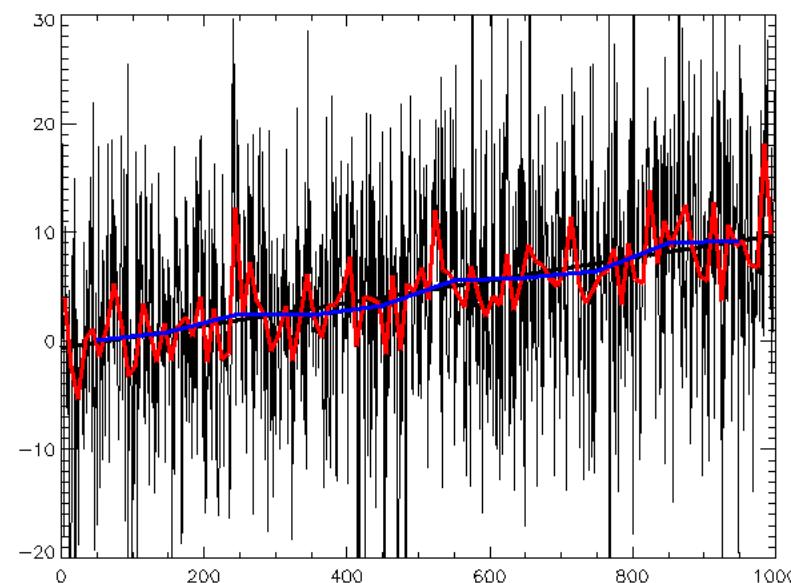
Outline

- Introduction
- Time-domain methods
 - Autoregressive integrated moving average (ARIMA)
 - Convolutional analysis and CNNs
- Frequency-domain methods
 - Spectral density analysis
 - Wavelet analysis
- Transformer for time series
- Conclusion

1. Introduction

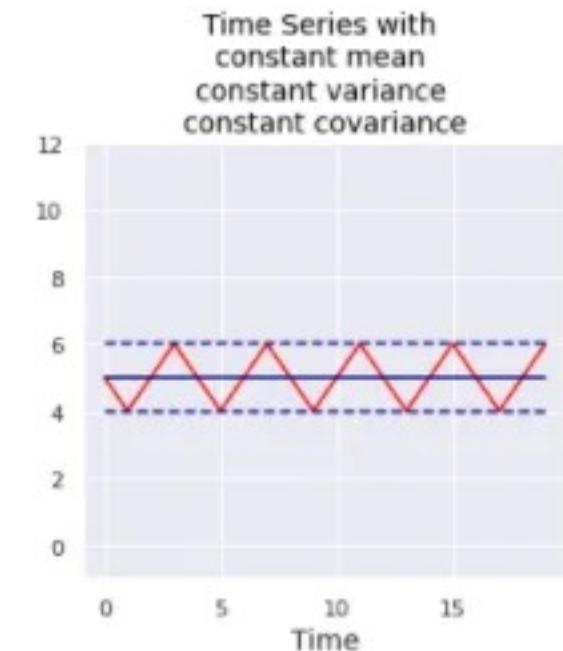
Time Series

- A sequence of data points read at equally spaced points of time
 - E.g. daily stock price, monthly rice price, heights of ocean tides, audio signals, counts of celestial meteriorites, and activity of tectonic plates
- Time series forecasting
 - Predicting future values based on previously observed values
 - **Stochastic process:** observations close together in time are more closely related than those further apart



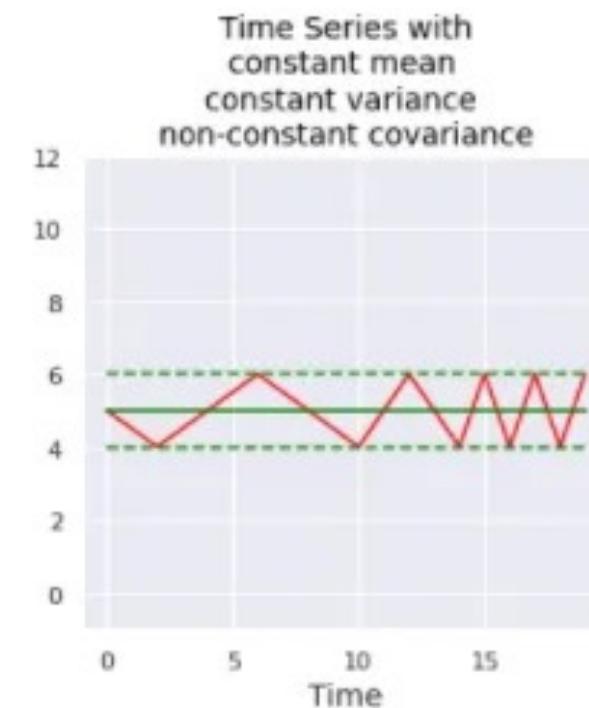
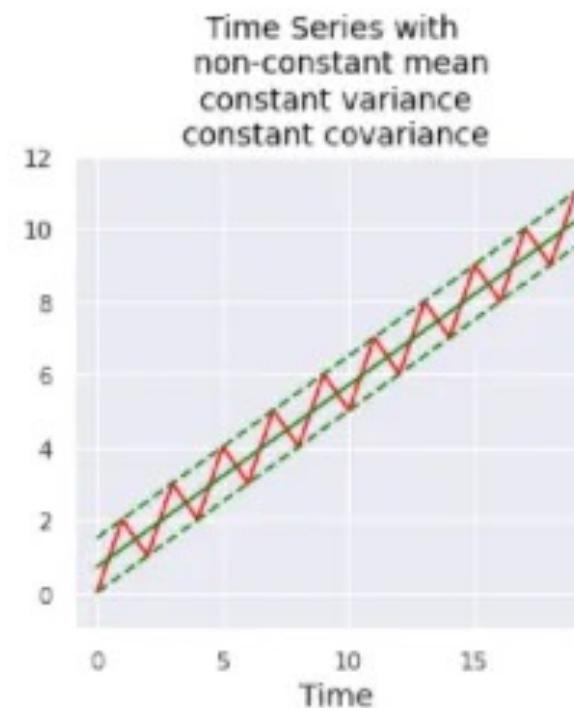
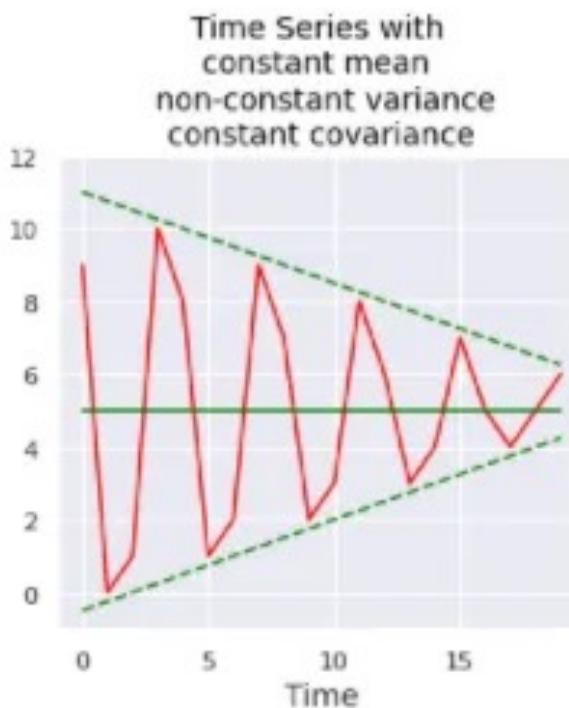
Stationary vs. Non-Stationary

- **Stationary time series** is one whose statistical properties do not change over time
 - **Constant mean:** average value remains the same throughout the time series
 - **Constant variance:** spread of data points do not change
 - **Constant autocovariance:** relationship between past and present values depends on the lag, not the time
- Non-stationary time series is one whose any of these properties change over time



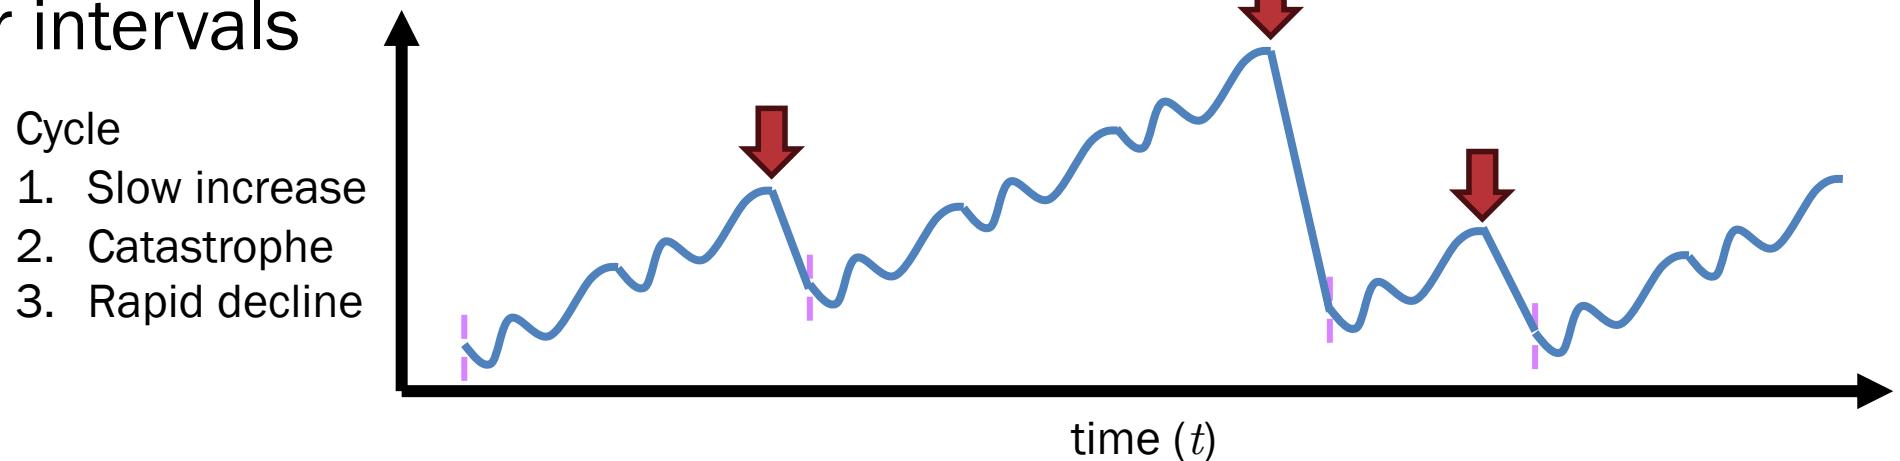
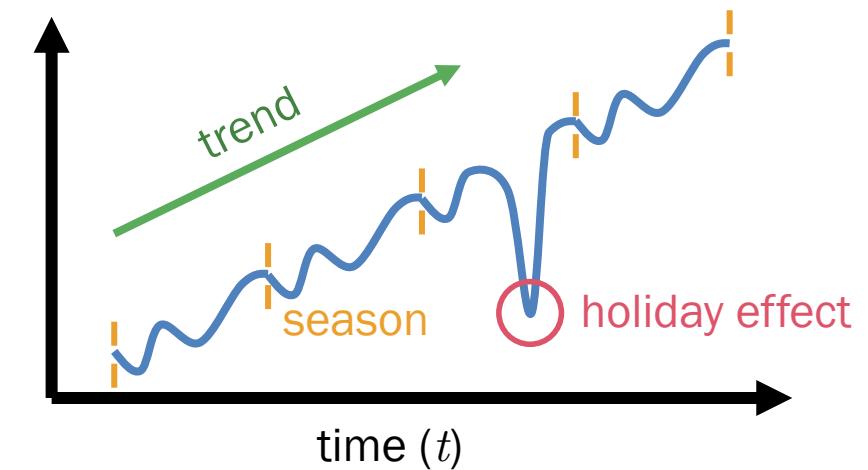
Stationary vs. Non-Stationary

- **Non-stationary time series** is one whose mean, variance (spread), or autocovariance (lag) change over time



Patterns in Time Series

- **Trend:** general direction over time
- **Seasonality:** repetitive patterns that occur at regular predictable intervals
- **Holiday effects:** irregular patterns caused by special calendar events
- **Cycle:** long-term repetitive patterns that occur at irregular intervals



Assumptions about Time Series

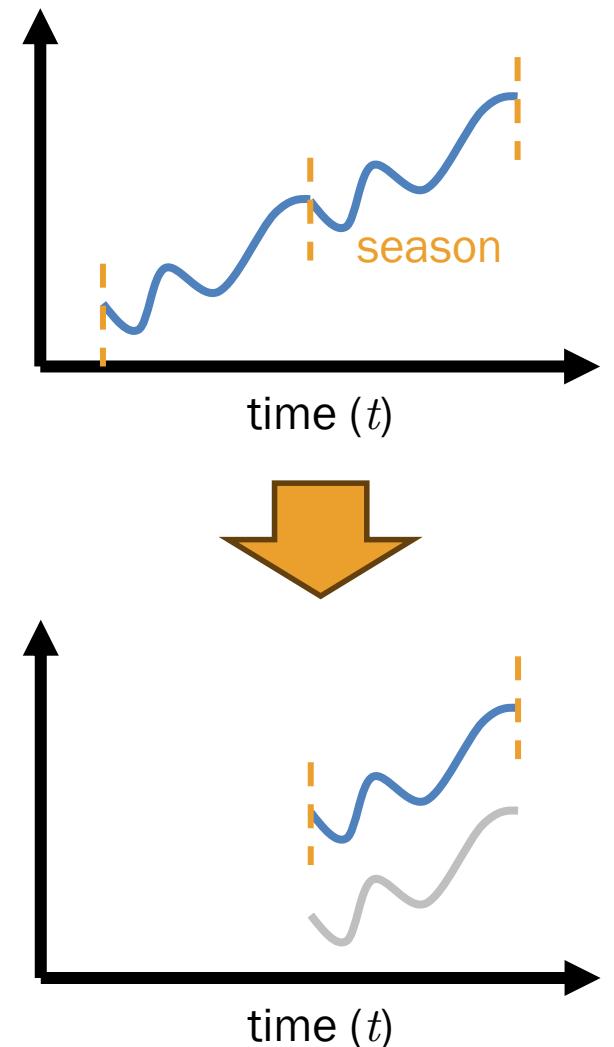
- Time domain
 - **A1:** Stochastic process (observations closer in time are more closely related)
 - **A2:** Combination of temporal structures
- Frequency domain
 - **A3:** Combination of continuous waves
 - **A4:** Combination of wavelets (i.e. wavelike pieces)
- Sequence-to-sequence prediction
 - **A5:** Transformer-based models

2. Time-Domain Methods

2.1 ARIMA Model

ARIMA Model

- **Auto-Regressive Integrated Moving Average**
 - **Assumption:** The dataset is seasonal and the difference between seasons can be predicted by linear regression
 - **How:** Predict a future value across the season by linear regression of previous cross-seasonal differences and adjust the error by linear regression of previous errors
- Three parameters of ARIMA(p, d, q)
 - Season duration: d timesteps
 - No. cross-seasonal differences: p timesteps
 - No. previous errors: q timesteps



ARIMA(p, d, q)

- **Step 1:** Integrate the season of length d

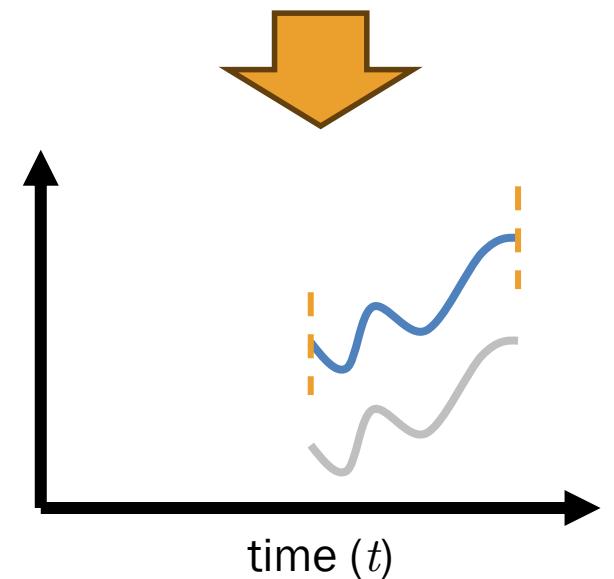
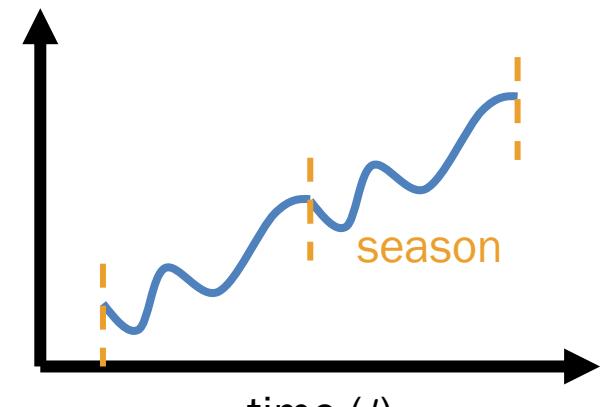
$$x'_t = x_t - x_{t-d}$$

for every point x_t in the time series

- **Example:**

Season length $d = 3$

t	x_t
1	10
2	15
3	20
4	25
5	28
6	38



ARIMA(p, d, q)

- **Step 1:** Integrate the season of length d

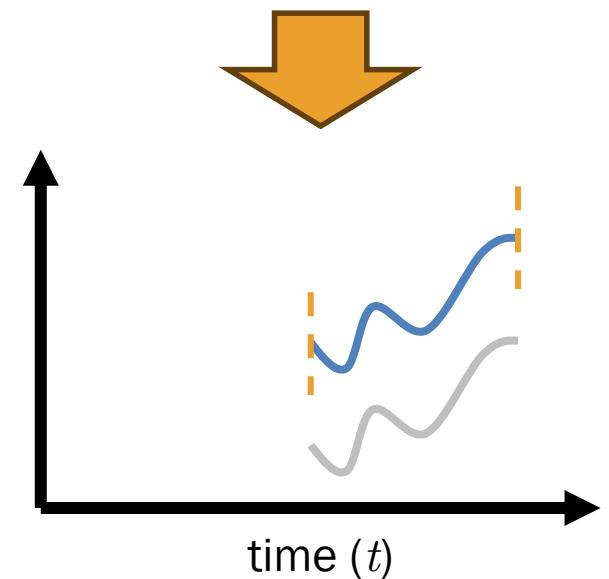
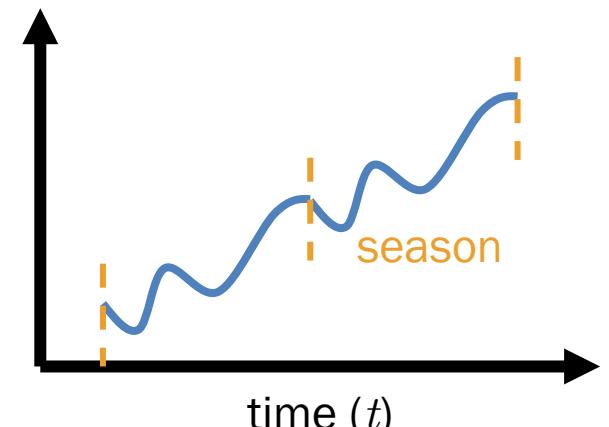
$$x'_t = x_t - x_{t-d}$$

for every point x_t in the time series

- **Example:**

Season length $d = 3$

t	x_t	x_{t-d}	x'_t
1	10	—	—
2	15	—	—
3	20	—	—
4	25	10	15
5	28	15	13
6	38	20	18



ARIMA(p, d, q)

- **Step 1:** Integrate the season of length d

$$x'_t = x_t - x_{t-d}$$

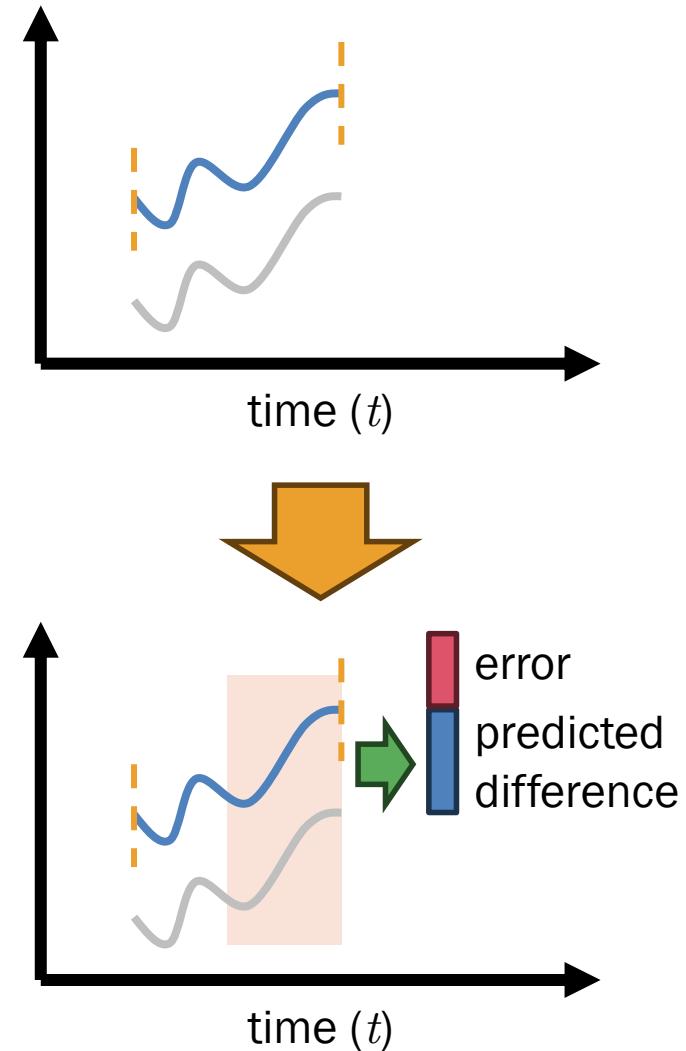
- **Step 2:** Predict the current difference with a linear regression of p previous differences

$$x'_t = \left[\sum_{k=1}^p \phi_k x'_{t-k} \right] + e_t$$

predicted difference
error

 $= x_t - (x_{t-d} + \text{diff})$

The term ‘autoregressive’ means taking the recent outputs as inputs for the next computation



ARIMA(p, d, q)

- **Step 1:** Integrate the season of length d

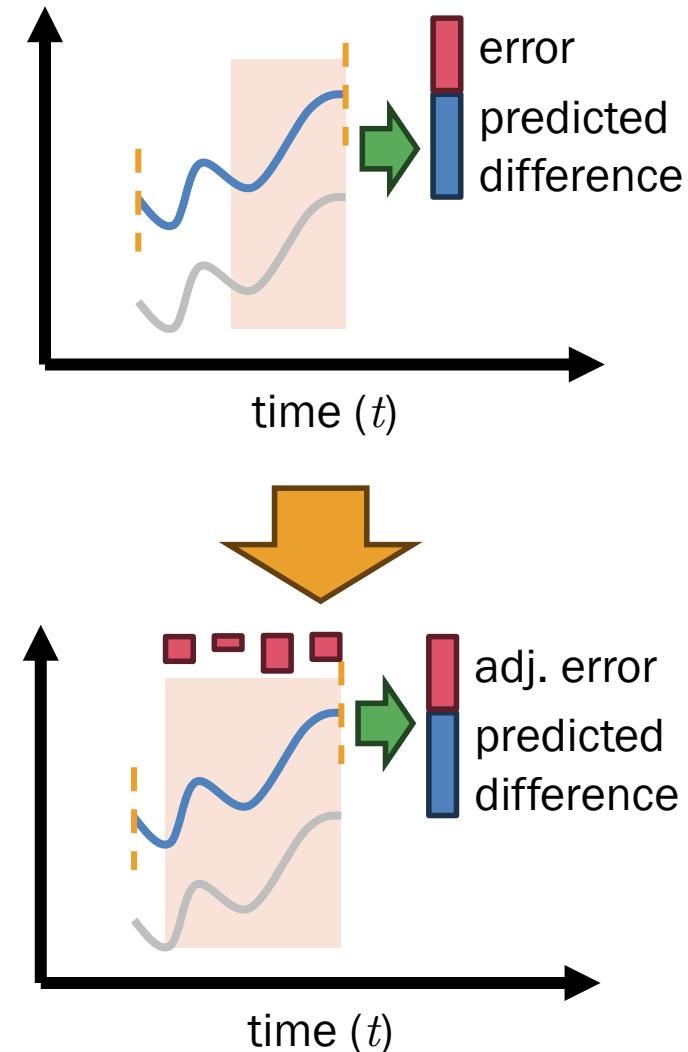
$$x'_t = x_t - x_{t-d}$$

- **Step 2:** Predict the current difference with a linear regression of p previous differences

$$x'_t = \left[\sum_{k=1}^p \phi_k x'_{t-k} \right] + e_t$$

- **Step 3:** Adjust the prediction error with a linear regression of q previous errors

$$x'_t = \left[\sum_{k=1}^p \phi_k x'_{t-k} \right] + e_t + \left[\sum_{j=1}^q \theta_j e_{t-j} \right]$$



Training Algorithm of ARIMA(p, d, q)

- Suppose each data point x_t is in the training set
 - Compute the differences x'_t for each data point

$$x'_t = x_t - x_{t-d}$$

- Estimate parameters φ_k with MSE of all e_t

$$e_t = x'_t - \left[\sum_{k=1}^p \phi_k x'_{t-k} \right]$$

- Estimate parameters θ_k with MSE of all e_t with fixed φ_k

$$e_t = x'_t - \left[\sum_{k=1}^p \phi_k x'_{t-k} \right] - \left[\sum_{j=1}^q \theta_j e_{t-j} \right]$$

$$\text{MSE} = \frac{1}{N} \sum_{t=1}^N e_t^2$$

Prediction of ARIMA(p,d,q)

- Suppose we want to predict future values x_{t+1} to x_{t+N}
 - We compute the differences and errors of timesteps $t+1$ to $t+N$

$$x'_t = \sum_{k=1}^p \phi_k x'_{t-k} + \sum_{j=1}^q \theta_j e_{t-j}$$

$$e_t = x'_t - \left[\sum_{k=1}^p \phi_k x'_{t-k} \right] - \left[\sum_{j=1}^q \theta_j e_{t-j} \right]$$

- We compute the future values from the differences

$$x_t = x'_t + x_{t-d}$$

Incomplete Time Series

- Interpolation techniques

- **Newton's interpolation:** each a_i is computed by divided differences

$$N(x) = a_0 + a_1(x - x_0) + a_2(x - x_0)(x - x_1) + \dots + a_n \prod_{k=0}^{n-1} (x - x_k)$$

- **Cubic spline interpolation:** cubic curve

$$S_i(x) = a_i + b_i(x - x_i) + c_i(x - x_i)^2 + d_i(x - x_i)^3$$

- **Chebyshev's interpolation:** sinusoidal seasonality

$$x_i = \frac{a+b}{2} + \frac{b-a}{2} \cos\left(\frac{2i+1}{2n}\pi\right)$$

- **Radial basis function interpolation:** spread around the means x_i

$$R(x) = \sum_{i=1}^n \lambda_i \cdot \phi(|x - x_i|)$$

Evaluation of Prediction Models

- **Mean absolute error (MAE):** average absolute difference between prediction and gold standard
- **Root mean squared error (RMSE):** square root of mean squared difference between prediction and gold standard
- **Mean absolute percentage error (MAPE):** average percentage difference between prediction and gold standard (w.r.t. gold standard)
- **Forecast bias:** average bias (prediction – gold standard)

$$\text{MAE} = \frac{1}{N} \sum_{k=1}^N |y_k - \hat{y}_k|$$

$$\text{RMSE} = \sqrt{\frac{1}{N} \sum_{k=1}^N (y_k - \hat{y}_k)^2}$$

$$\text{MAPE} = \frac{1}{N} \sum_{k=1}^N \left| \frac{y_k - \hat{y}_k}{y_k} \right|$$

$$\text{bias} = \frac{1}{N} \sum_{k=1}^N (y_k - \hat{y}_k)$$

2.2 Convolutional Analysis and CNNs

Convolution *

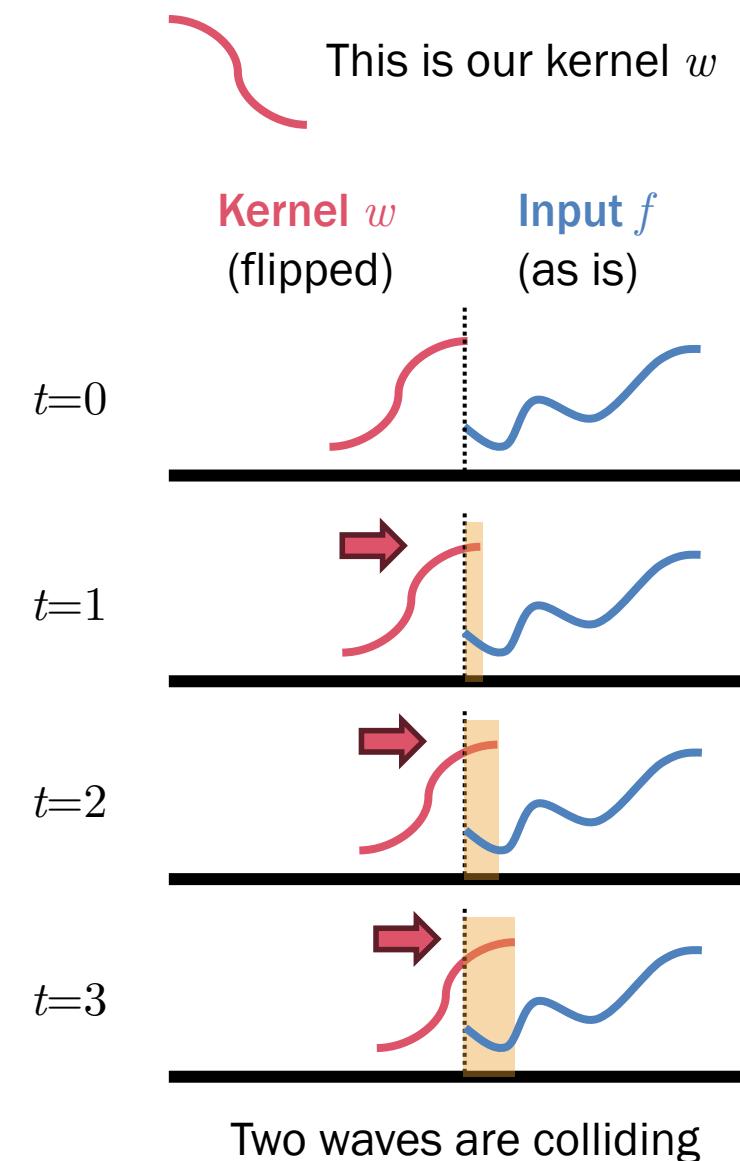
- Combining two functions to produce a third one that represents how one amplifies the other in each step of their overlapping
 - Continuous convolution**

$$[f * w](t) = \int_{-\infty}^{\infty} f(\tau) \cdot w(t - \tau) d\tau$$

- Discrete convolution**

$$[f * w](t) = \sum_{k=-\infty}^{\infty} f[k] \cdot w[t - k]$$

- where f is an input signal and w is a kernel (i.e. pattern filter)



Cross-Correlation \odot

- The intuitive counterpart of convolution, where the kernel function is not flipped

- Continuous cross-correlation**

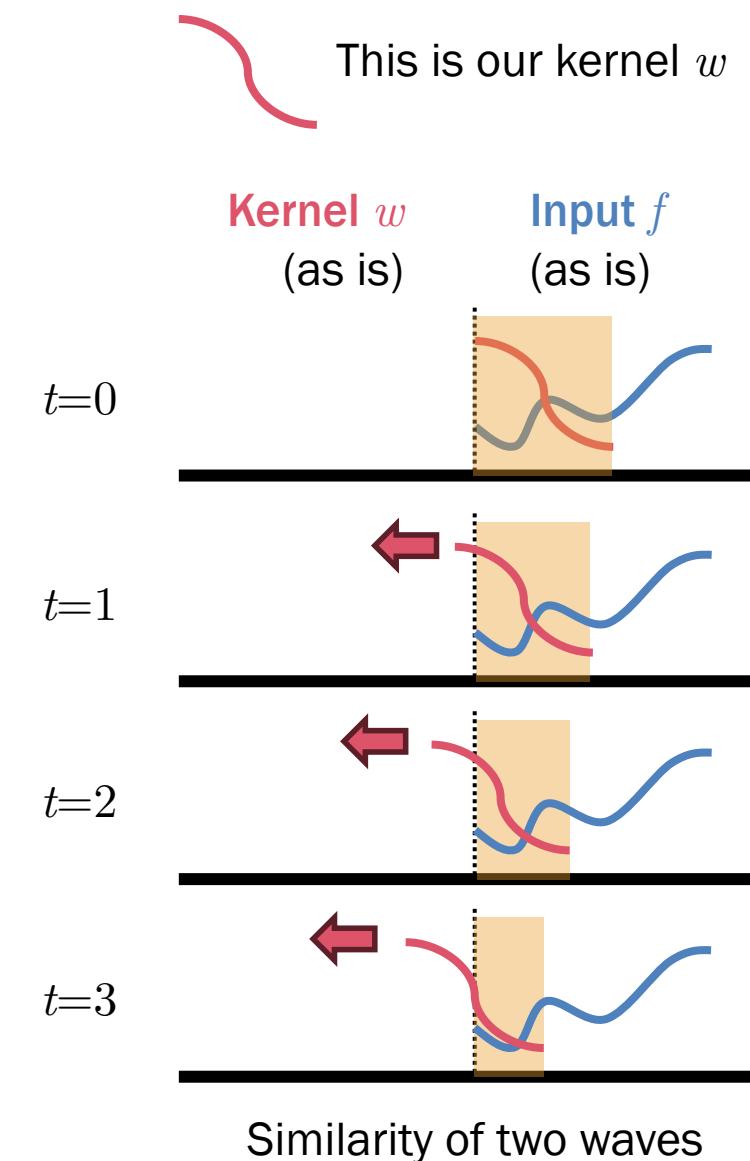
$$[f \odot w](t) = \int_{-\infty}^{\infty} f(\tau) \cdot w(t + \tau)^* d\tau$$

* is the complex conjugate
e.g. $(3 + 4i)^* = 3 - 4i$

- Discrete cross-correlation**

$$[f \odot w](t) = \sum_{k=-\infty}^{\infty} f[k] \cdot w[t + k]^*$$

- We will say '**convolution**' to refer to cross-correlation for ease of understanding



Convolutional Filtering



Find
(kernel)



Convolutional Filtering



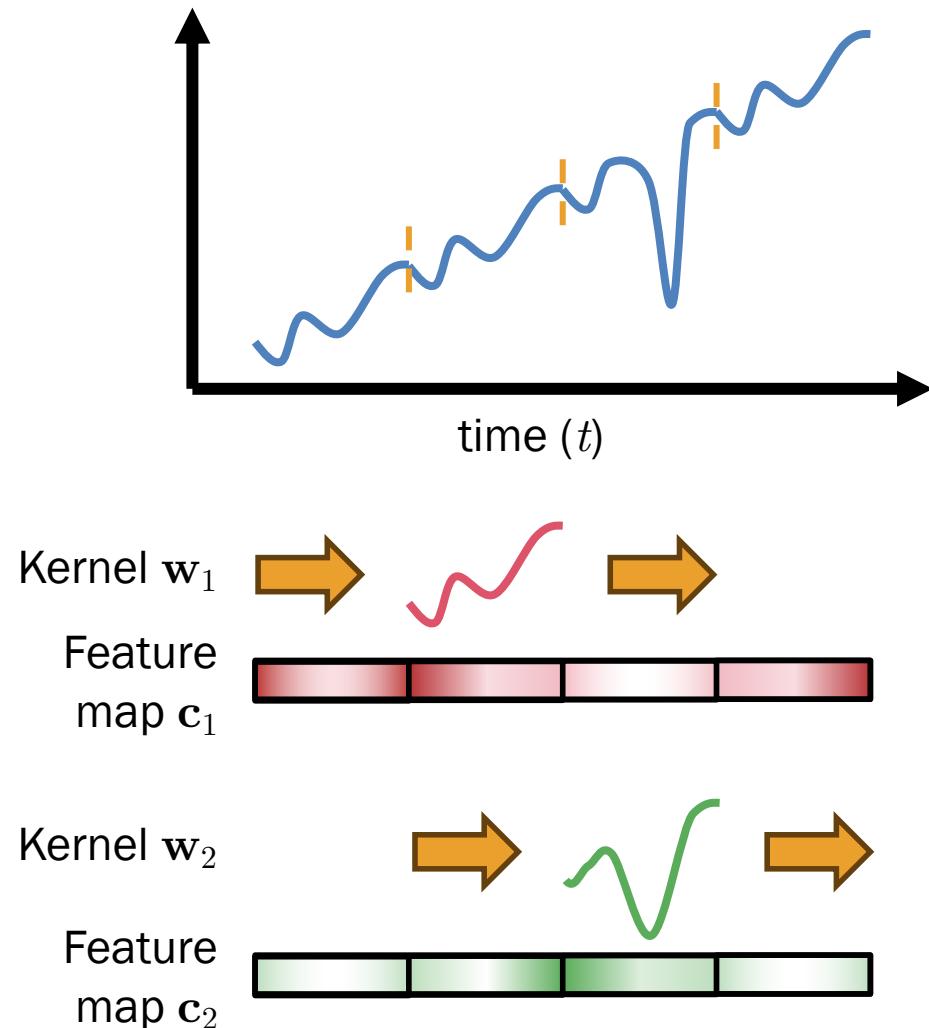
Find
(kernel)



- Matched areas are amplified
- Objects are detected via convolution
- We have extracted the **local features**

Convolution in Time Series

- Pattern matching with kernels
 - **Assumption:** Local features are detected by a series of peaks
 - Period of seasonality can be identified by the peaks of cross-correlations
 - Holiday effects are also identified by the irregularity in cross-correlation peaks
 - One time series may align (correlate) with a mixture of kernels



Convolutional Neural Networks

- Learn several kernels from the dataset and identify which kernels to be used at which time
 - Convolution layer with 3 kernels (width=10)

$$\mathbf{c}_i = \mathbf{x} \odot \mathbf{w}_i \quad \text{N.B. These kernels can learn any shape of trends}$$

where each \mathbf{w}_i is a vector of 10 parameters

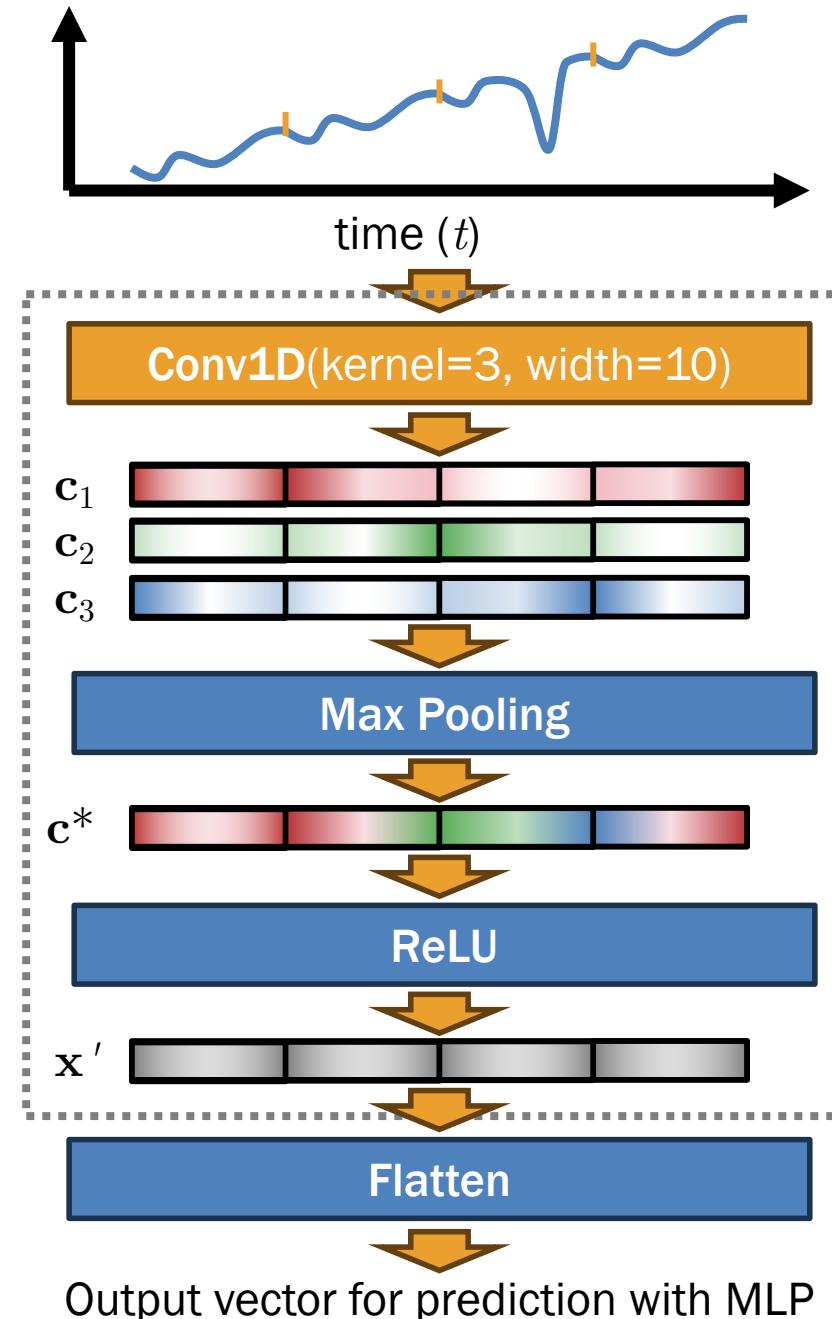
- Max-pooling for most prominent local features

$$\mathbf{c}^* = \max [\mathbf{c}_1 | \mathbf{c}_2 | \mathbf{c}_3]$$

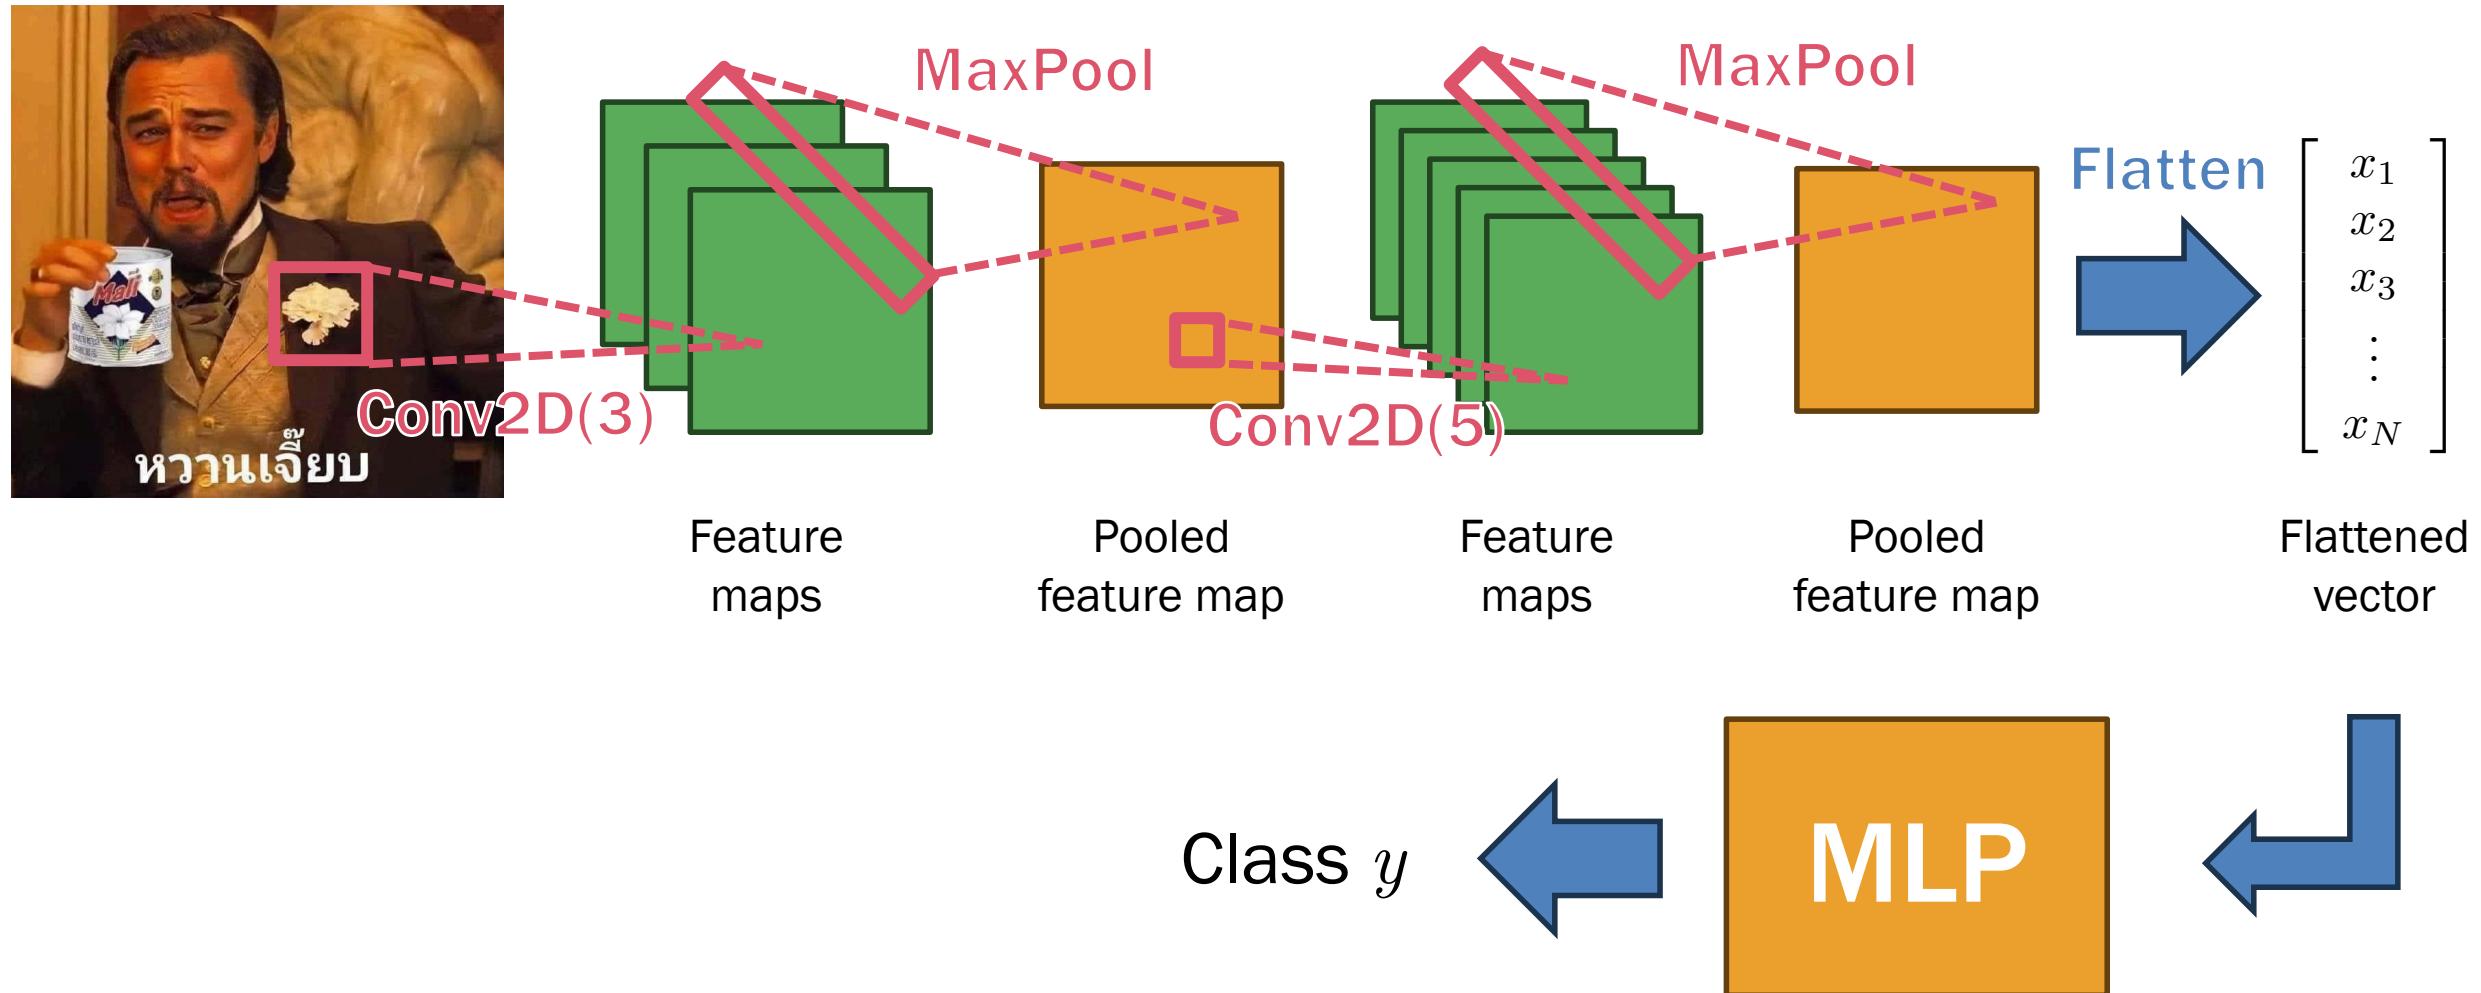
- Nonlinear function

$$\mathbf{x}' = \text{ReLU}(\mathbf{c}^*)$$

- The matrix is then flattened to become a vector



CNNs for Computer Vision

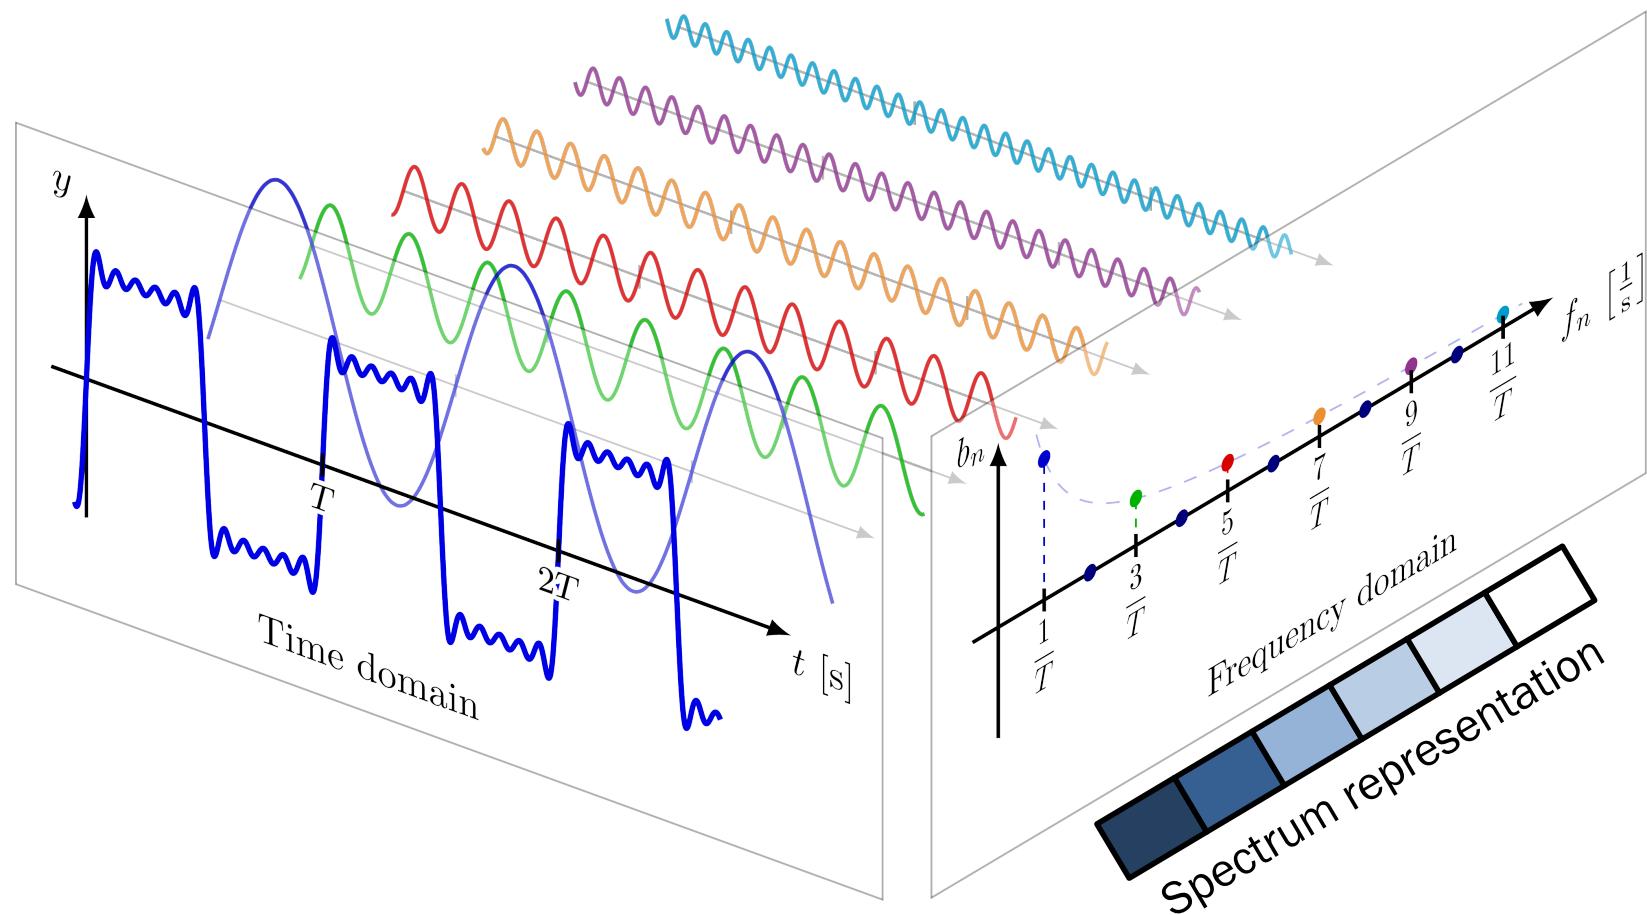


3. Frequency-Domain Methods

3.1 Spectral Density Analysis

Spectral Density

- Describing a time series in terms of power distribution according to frequency components
- Assumption:** Time series is a mixture of periodic sinusoid waves (seasonal patterns)
- Transforming a time series into **spectrum** (squared amplitude) in the frequency domain

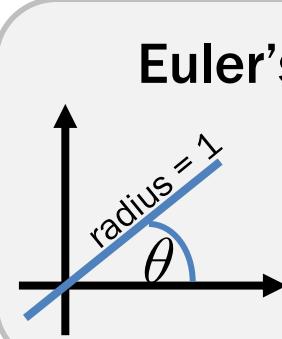


Source: <https://dibsmethodsmeetings.github.io/fourier-transforms/>

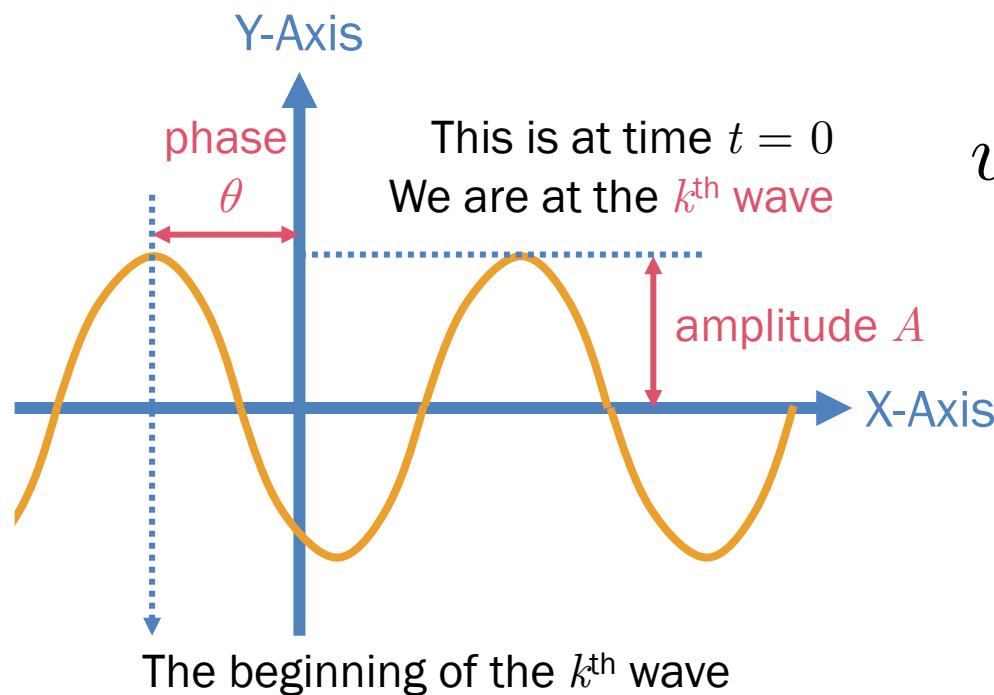
Wave as Complex Number

- A complex number $a + \hat{i}b$ consists of the real part a and the imaginary part b , where $\hat{i} = \sqrt{-1}$
- Wave can be represented as a complex number

Euler's formula



$$\exp[\hat{i}\theta] = \cos \theta + \hat{i} \sin \theta$$

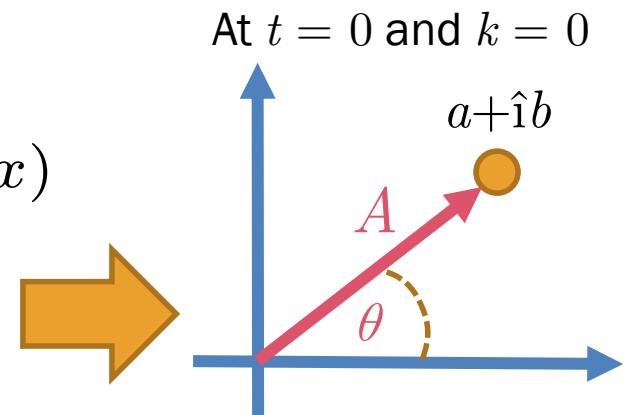


$$u(x, t) = A \cos(kx - \omega t + \theta)$$

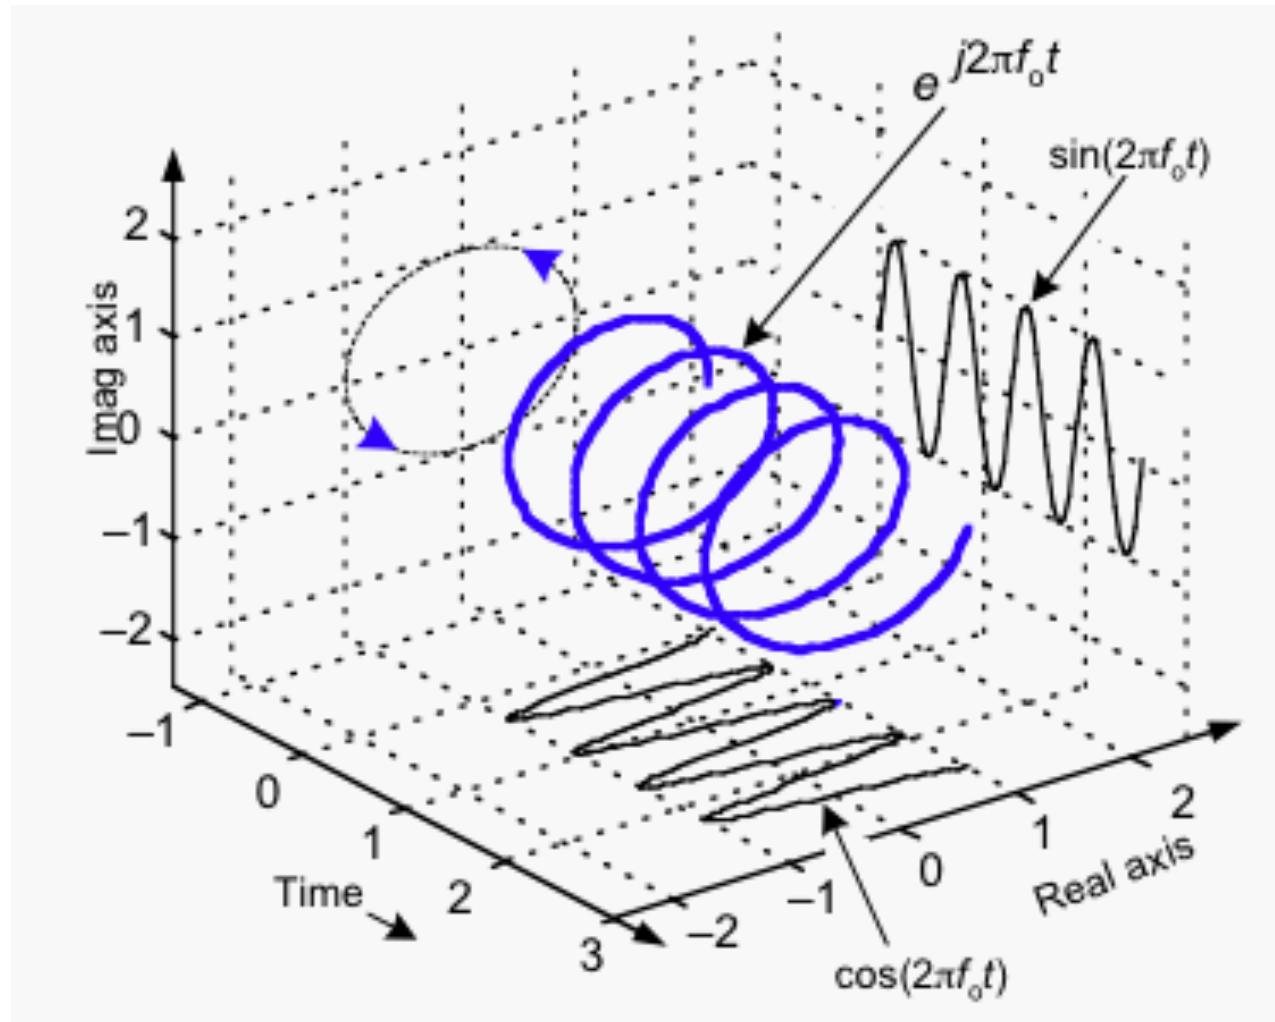


$$u(x, t) = u_0 e^{-\hat{i}(\omega t - kx)}$$

where $u_0 = A e^{\hat{i}\theta}$



Wave as 3D Spiral



- Wave can be seen as a counterclockwise spiral in 3D space, whose base plane are real and imaginary parts
$$\exp[\hat{i}\omega t] = \cos \omega t + \hat{i} \sin \omega t$$

real part imaginary part

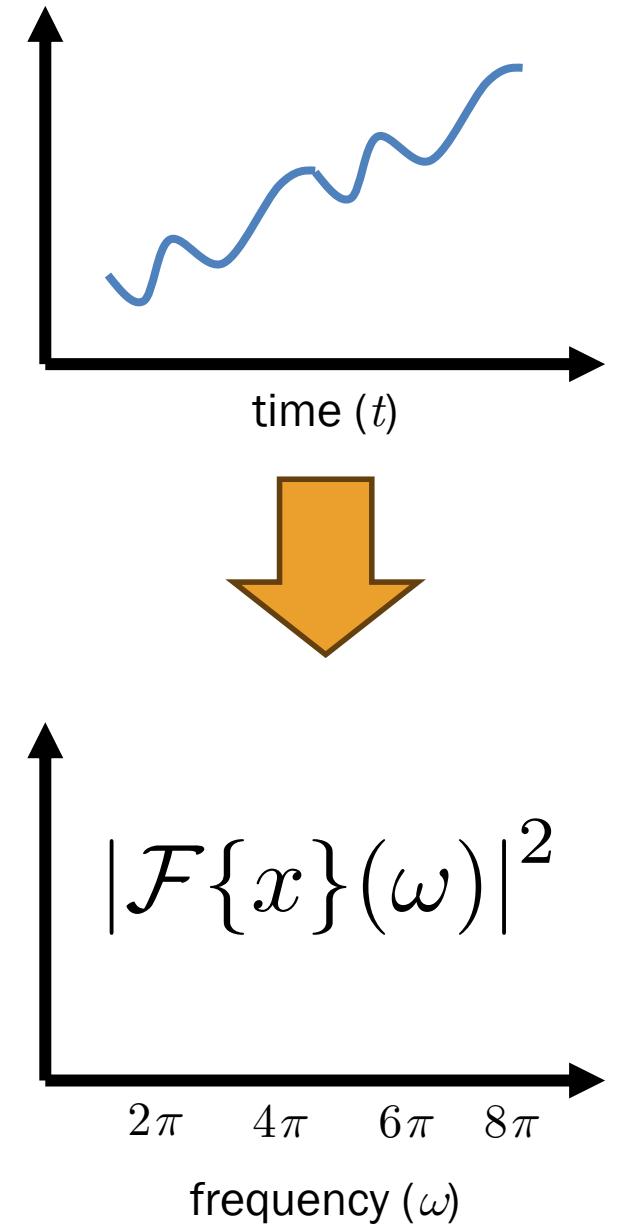
- Conjugate of a wave is therefore a clockwise spiral
$$\begin{aligned}\exp [\hat{i}\omega t]^* &= \exp [-\hat{i}\omega t] \\ &= \cos \omega t - \hat{i} \sin \omega t\end{aligned}$$

Fourier Transform

- Fourier transform of a time series $x(t)$ is a function of frequency ω that reflects how well $x(t)$ aligns with the wave of frequency ω

$$\begin{aligned}\mathcal{F}\{x\}(\omega) &= x(t) \odot \text{wave}(\omega, t) \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[\hat{i}\omega t]^* dt \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[-\hat{i}\omega t] dt\end{aligned}$$

- The frequency is usually a multiple of 2π (i.e. one round of a circle in radian)

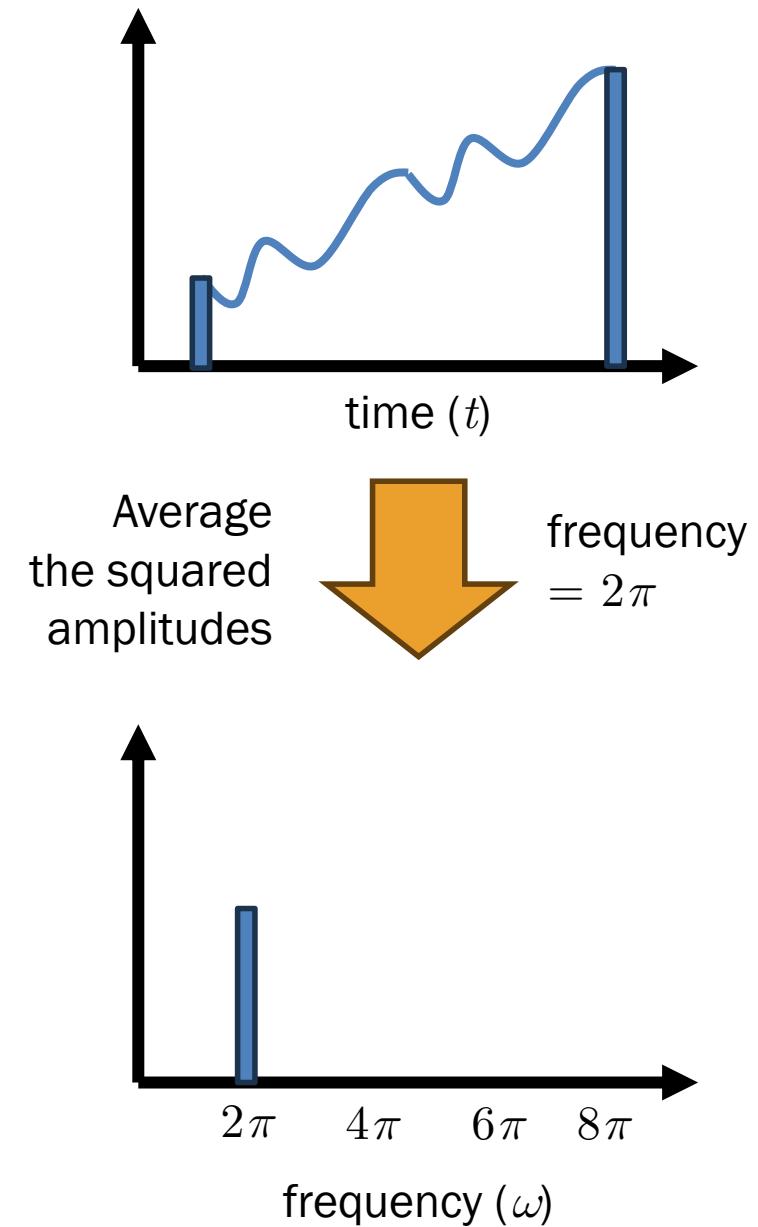


Fourier Transform

- Fourier transform of a time series $x(t)$ is a function of frequency ω that reflects how well $x(t)$ aligns with the wave of frequency ω

$$\begin{aligned}\mathcal{F}\{x\}(\omega) &= x(t) \odot \text{wave}(\omega, t) \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[\hat{i}\omega t]^* dt \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[-\hat{i}\omega t] dt\end{aligned}$$

- The frequency is usually a multiple of 2π (i.e. one round of a circle in radian)

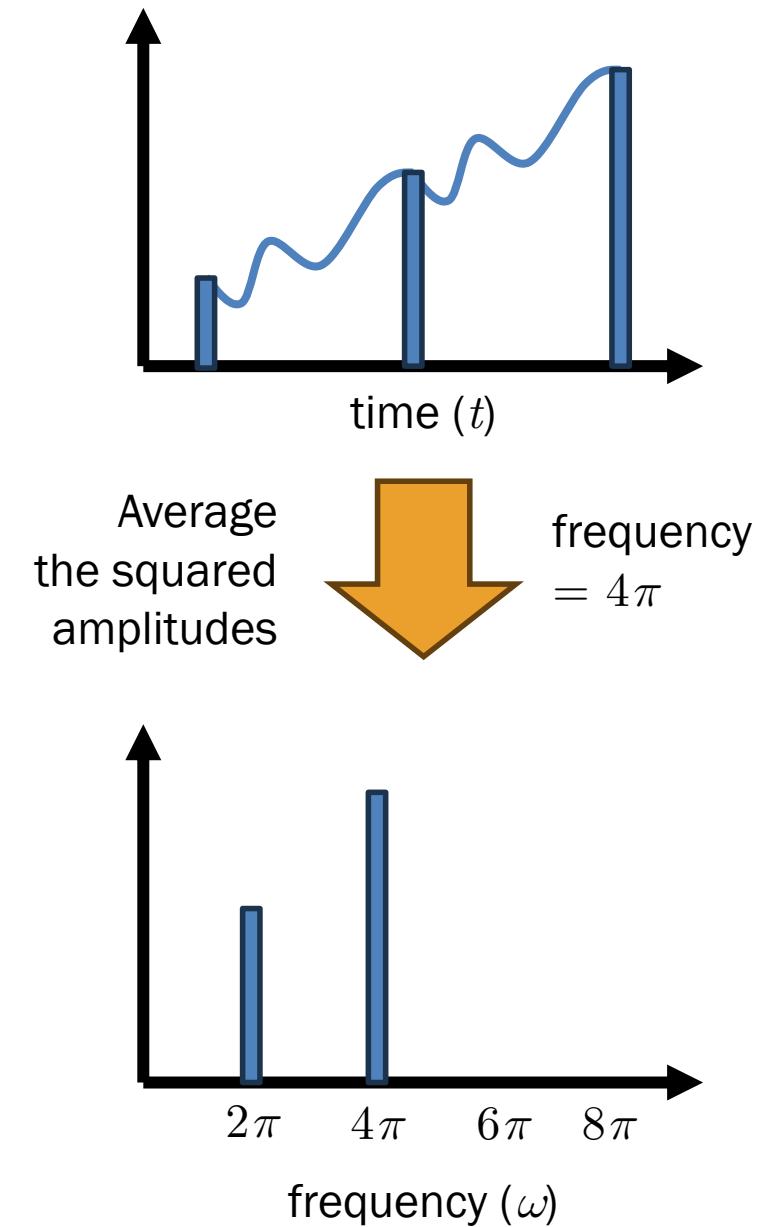


Fourier Transform

- Fourier transform of a time series $x(t)$ is a function of frequency ω that reflects how well $x(t)$ aligns with the wave of frequency ω

$$\begin{aligned}\mathcal{F}\{x\}(\omega) &= x(t) \odot \text{wave}(\omega, t) \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[\hat{i}\omega t]^* dt \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[-\hat{i}\omega t] dt\end{aligned}$$

- The frequency is usually a multiple of 2π (i.e. one round of a circle in radian)

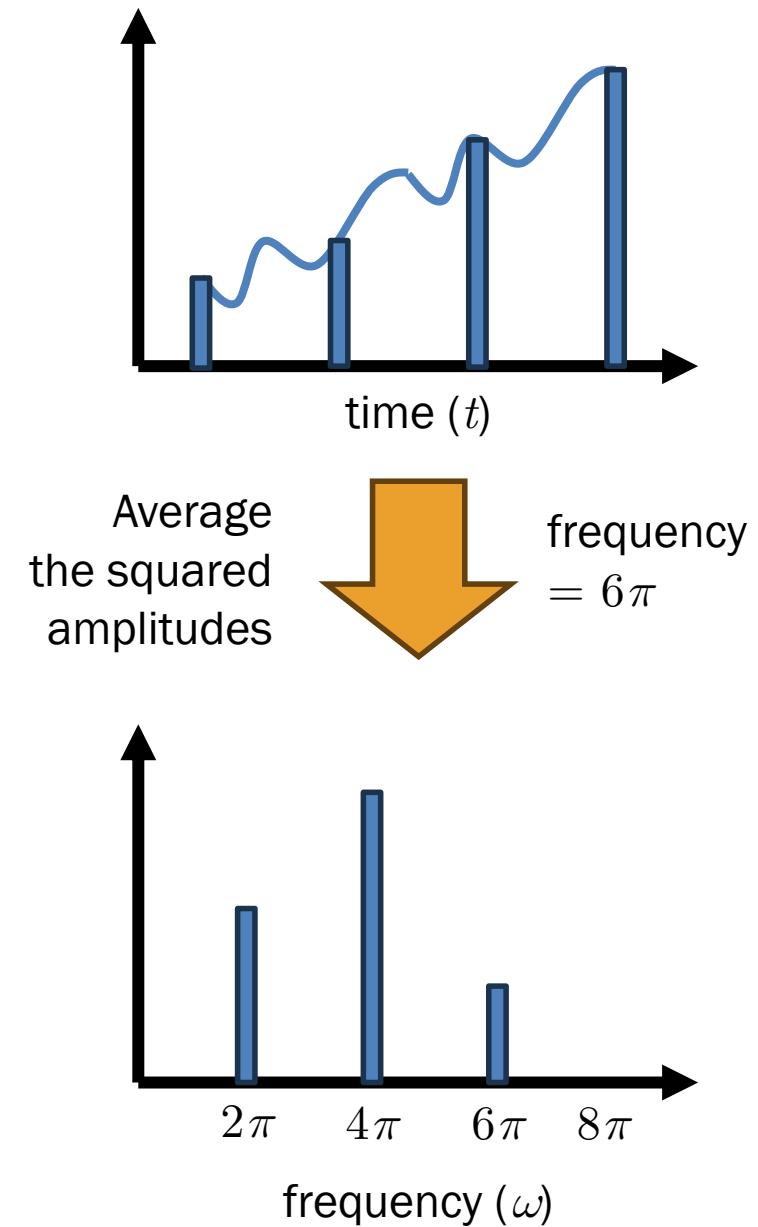


Fourier Transform

- Fourier transform of a time series $x(t)$ is a function of frequency ω that reflects how well $x(t)$ aligns with the wave of frequency ω

$$\begin{aligned}\mathcal{F}\{x\}(\omega) &= x(t) \odot \text{wave}(\omega, t) \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[\hat{i}\omega t]^* dt \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[-\hat{i}\omega t] dt\end{aligned}$$

- The frequency is usually a multiple of 2π (i.e. one round of a circle in radian)

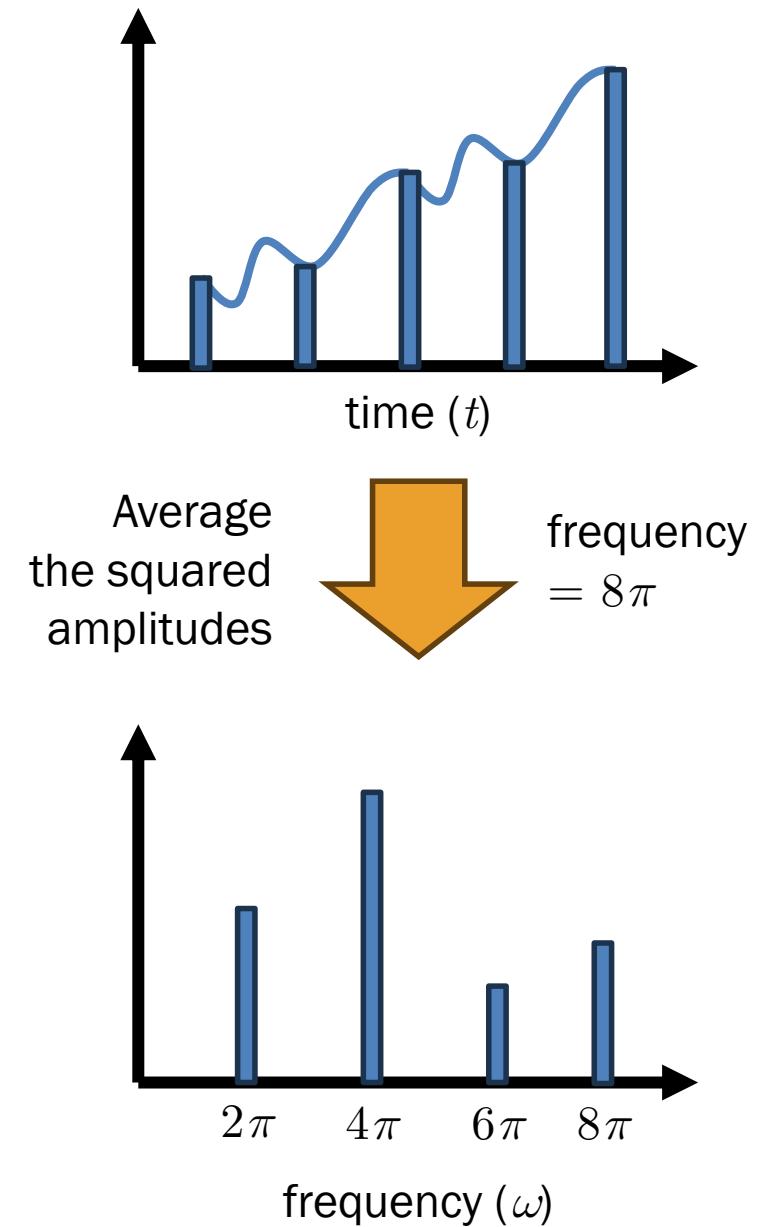


Fourier Transform

- Fourier transform of a time series $x(t)$ is a function of frequency ω that reflects how well $x(t)$ aligns with the wave of frequency ω

$$\begin{aligned}\mathcal{F}\{x\}(\omega) &= x(t) \odot \text{wave}(\omega, t) \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[\hat{i}\omega t]^* dt \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[-\hat{i}\omega t] dt\end{aligned}$$

- The frequency is usually a multiple of 2π (i.e. one round of a circle in radian)

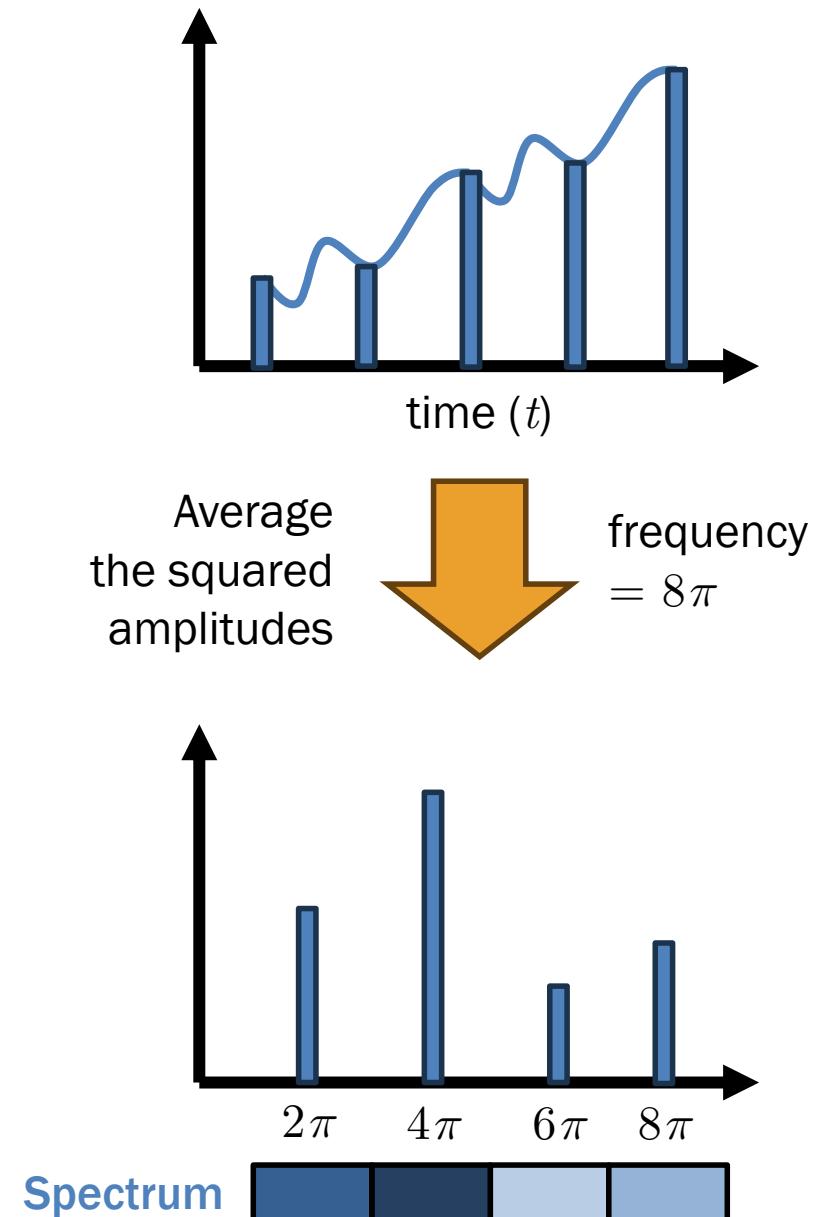


Fourier Transform

- Fourier transform of a time series $x(t)$ is a function of frequency ω that reflects how well $x(t)$ aligns with the wave of frequency ω

$$\begin{aligned}\mathcal{F}\{x\}(\omega) &= x(t) \odot \text{wave}(\omega, t) \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[\hat{i}\omega t]^* dt \\ &= \int_{-\infty}^{\infty} x(t) \cdot \exp[-\hat{i}\omega t] dt\end{aligned}$$

- The frequency is usually a multiple of 2π (i.e. one round of a circle in radian)



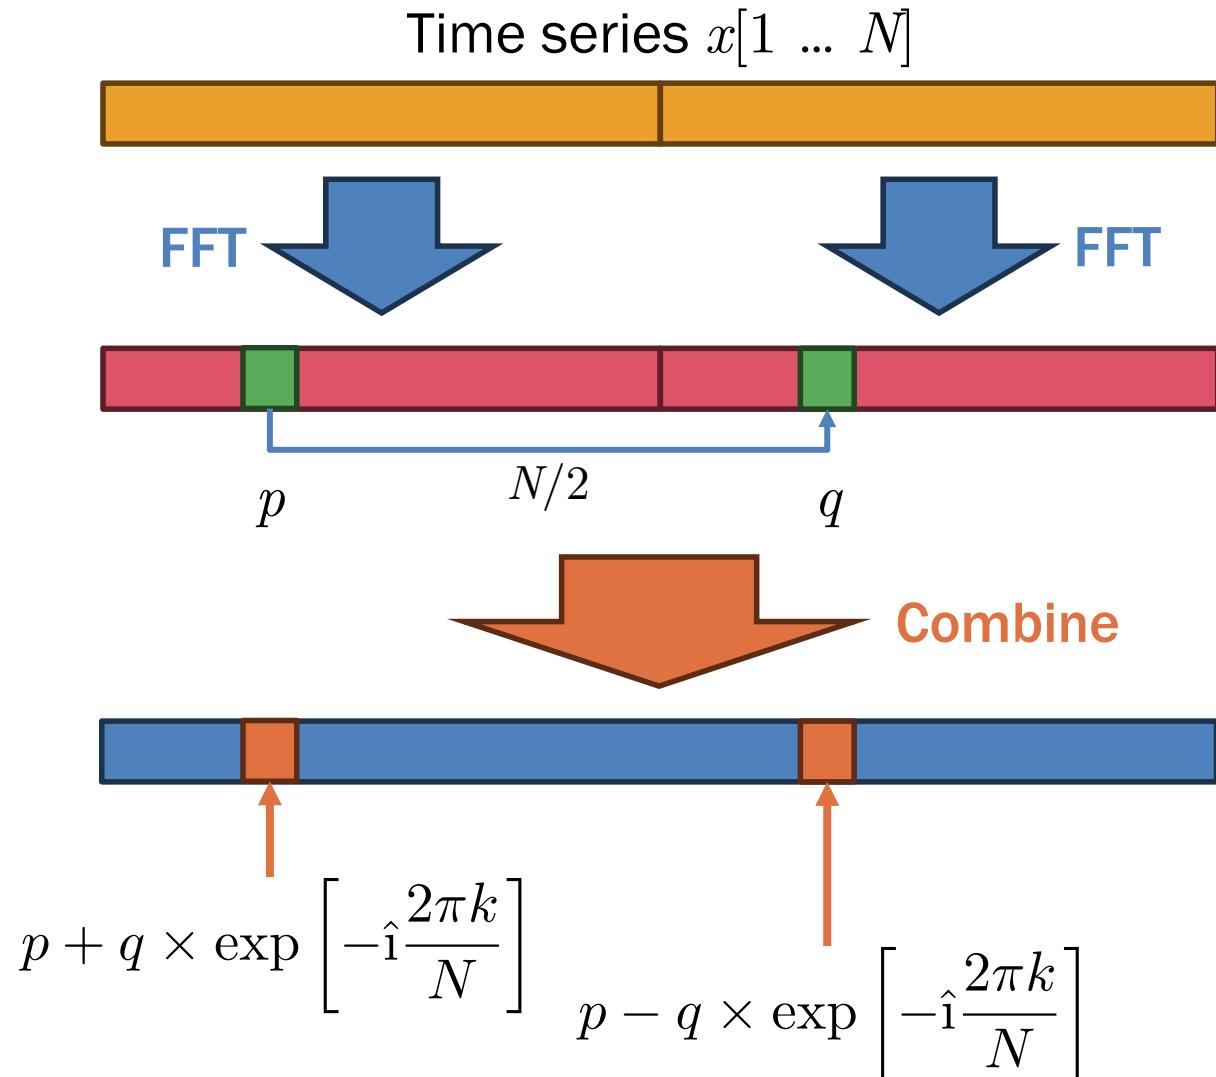
Fast Fourier Transform (Cooley & Tukey, 1965)

```
def fft(x: array, N: length, s: stride):  
    result := create_array(size=N)  
  
    if N == 1:  
        result[0] = x[0]  
    else:  
        result[0 : N/2] := fft(x, N/2, 2*s)          # x[0], x[2s], x[4s], ...  
        result[N/2 : N] := fft(x[s:], N/2, 2*s)      # x[s], x[3s], x[5s], ...  
        for k in range(N/2):  
            p := result[k]  
            q := result[k + N/2] * exp(-2 * π * i * k/N)  
            result[k] := p + q  
            result[k + N/2] := p - q  
  
    return result
```

$O(n \log n)$
time
complexity

HOW TO RUN ==> `fft(x, N, s=1)` ==> Then compute the squared magnitude of each element

Fast Fourier Transform (Cooley & Tukey, 1965)



- Base case

$$\mathcal{F}(x[k]) = x[k]$$

- Recursive case

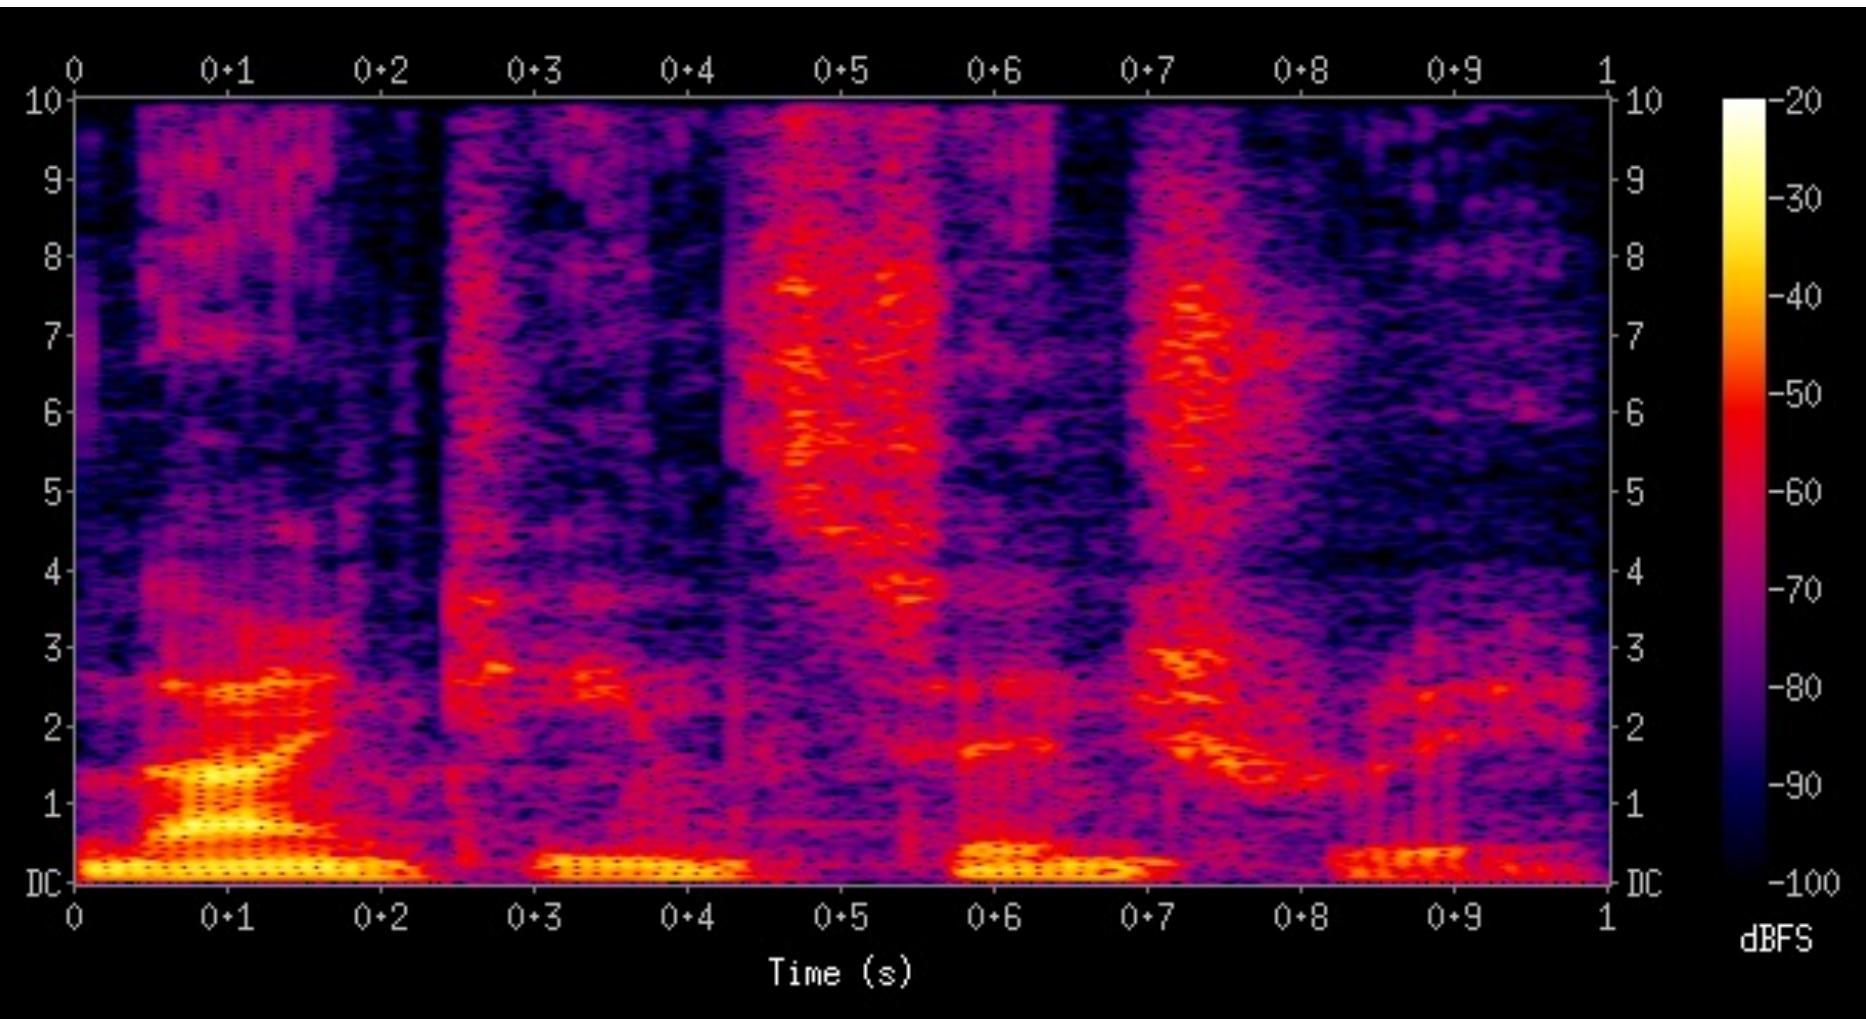
$$p = \mathcal{F}[x[k]]$$

$$q = \mathcal{F}[x[k + \frac{N}{2}]]$$

$$\mathcal{F}[x[k]] = p + q \times \exp\left[-i\frac{2\pi k}{N}\right]$$

$$\mathcal{F}[x[k + \frac{N}{2}]] = p - q \times \exp\left[-i\frac{2\pi k}{N}\right]$$

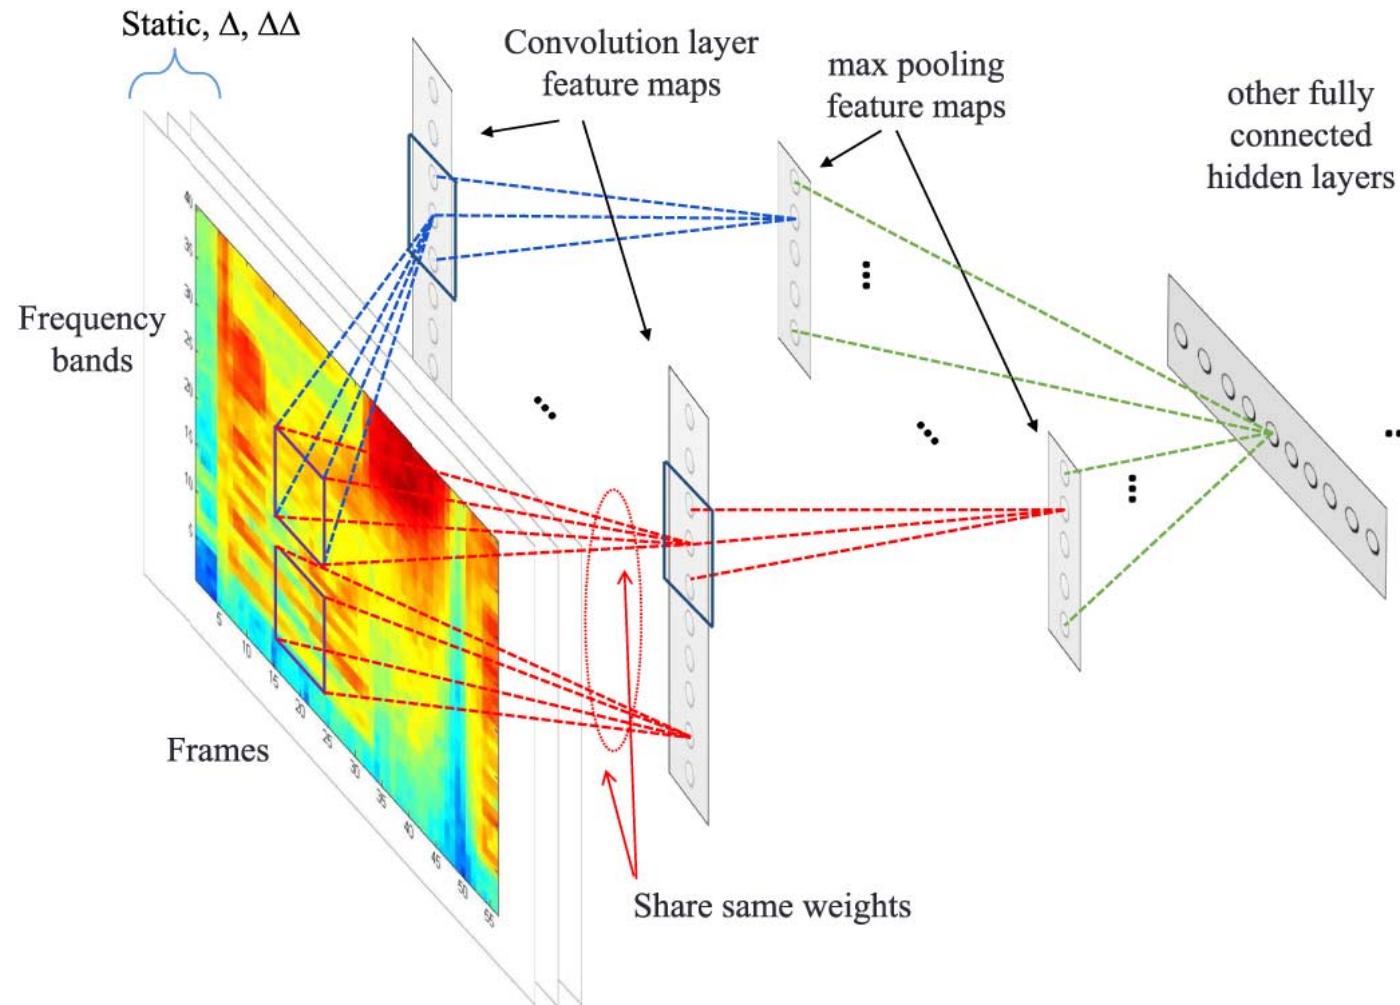
Spectrogram



- Heatmap is used for visual representation of the spectrum of frequency
- Time series is tokenized into equal chunks (w.r.t. window size) and analyzed with FFT
- Seasonality and holiday effects are present
- Good for stationary signals e.g. long-term climate data

Source: <https://en.wikipedia.org/wiki/Spectrogram>

Asynchronous Speech Recognition



- Spectrogram is used as an input picture for CNN-based models
- Local features are changes in specific frequency ranges
- Conv2D is usually employed

3.2 Wavelet Analysis

Wavelet Transform

- Wavelet transform of $x(t)$ is a function of frequency a and position b that reflects how well $x(t)$ aligns with the wavelet $\Psi(a, b, t)$

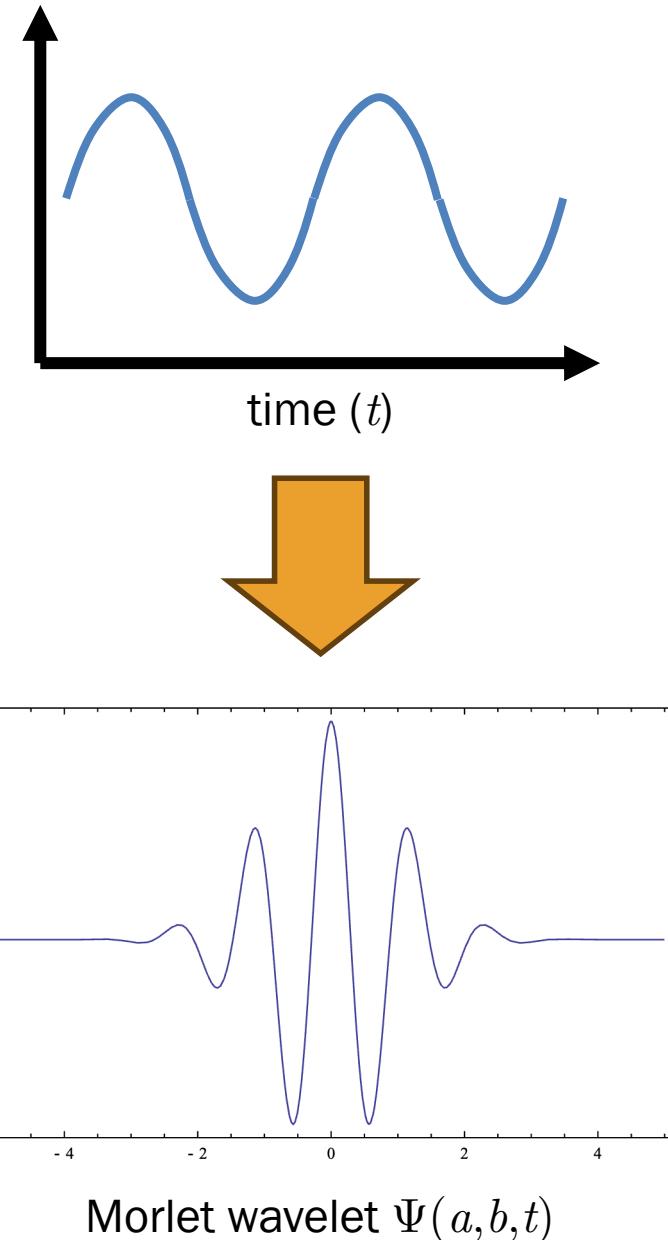
$$\mathcal{W}\{x\}(a, b) = x(t) \odot \Psi(a, b, t)$$

$$= \int_{-\infty}^{\infty} x(t) \cdot \cos a(t - b) \exp \left[(t - b)^2 \right] dt$$

where Morlet wavelet (1984) is used here:

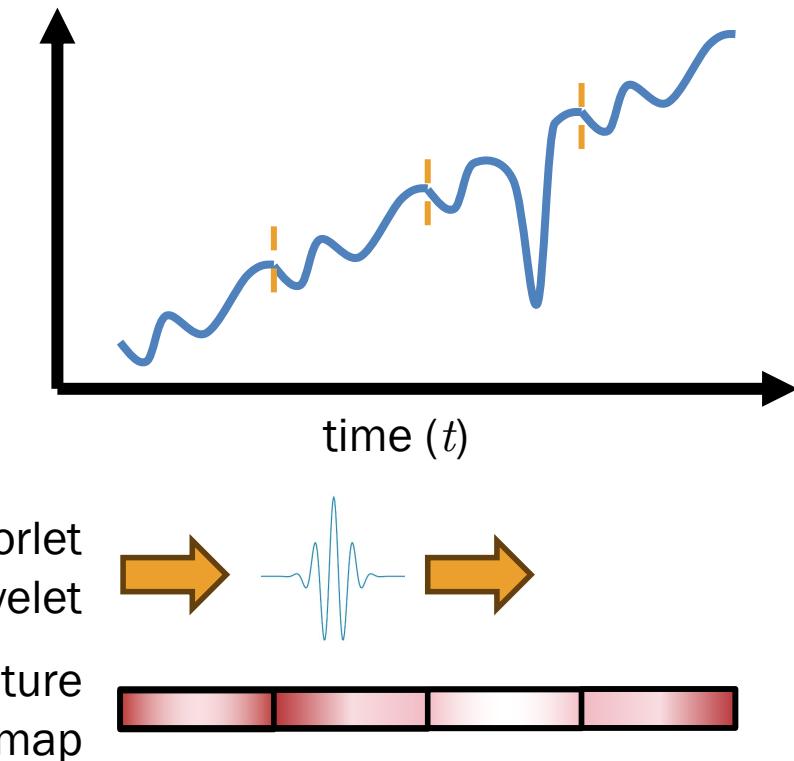
$$\Psi(a, b, t) = \cos a(t - b) \exp \left[(t - b)^2 \right]$$

- In practice, we need a low-pass filter $\Phi(t)$ to eliminate the inherent noise



Wavelet Transform in Time Series

- Pattern matching with wavelet
 - **Assumption:** Local features are detected by a series of peaks
 - Period of seasonality can be identified by the peaks of cross-correlations
 - Holiday effects are also identified by the irregularity in cross-correlation peaks
 - One time series may align (correlate) with a mixture of wavelets



Fast Wavelet Transform (Mallat, 1989)

```
def fwt(x: array, N: length, h: low-pass filter, g: wavelet, L: resolution level)
    result := []
    arr := x

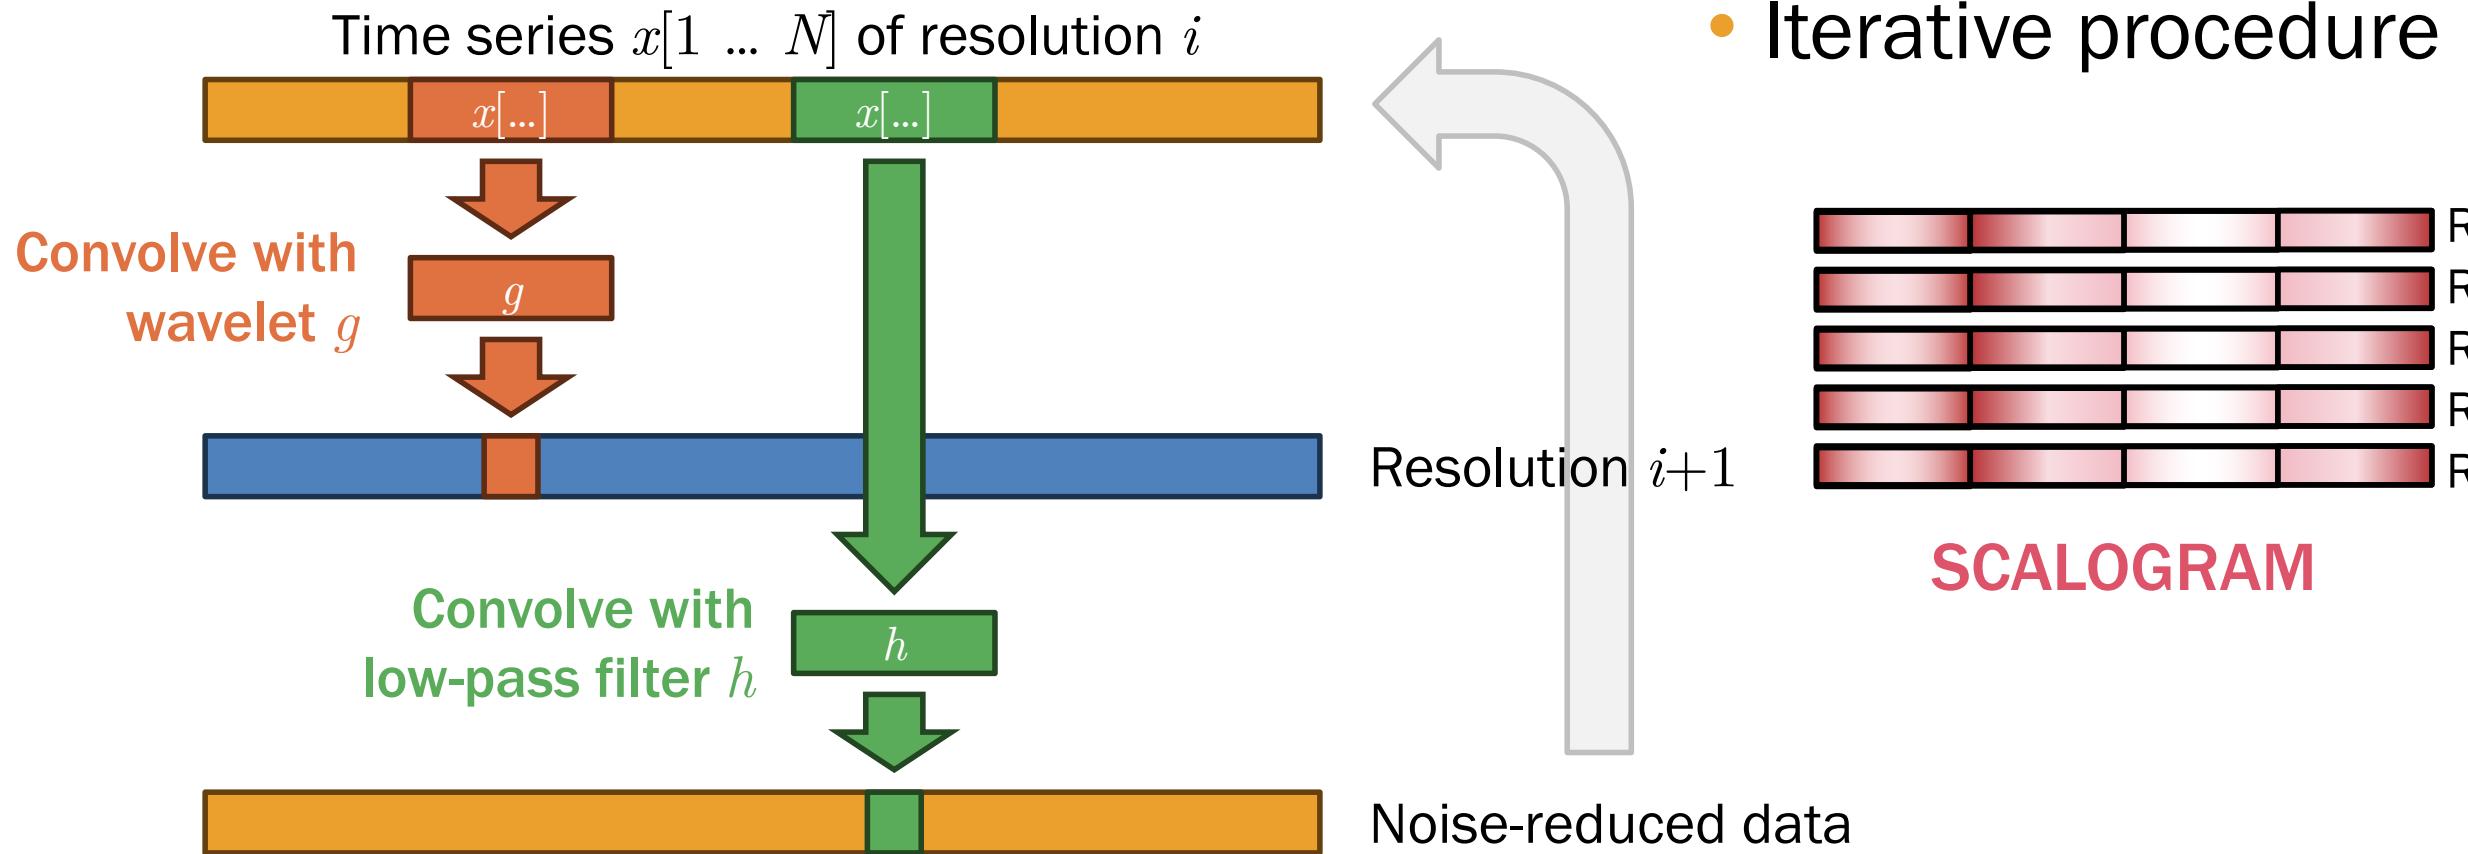
    for i = 1 to L:      # Resolution i
        row = []
        for k = 0 to N/2 - 1:
            coeff := sum(g[j] * arr[2*k - j] for j = 0 to len(g))
            row.append(coeff)
        result.insert(0, row)

        arr_next := []
        for k = 0 to N/2 - 1:
            coeff := sum(h[j] * arr[2*k - j] for j = 0 to len(h))
            arr_next.append(coeff)
        arr := arr_next

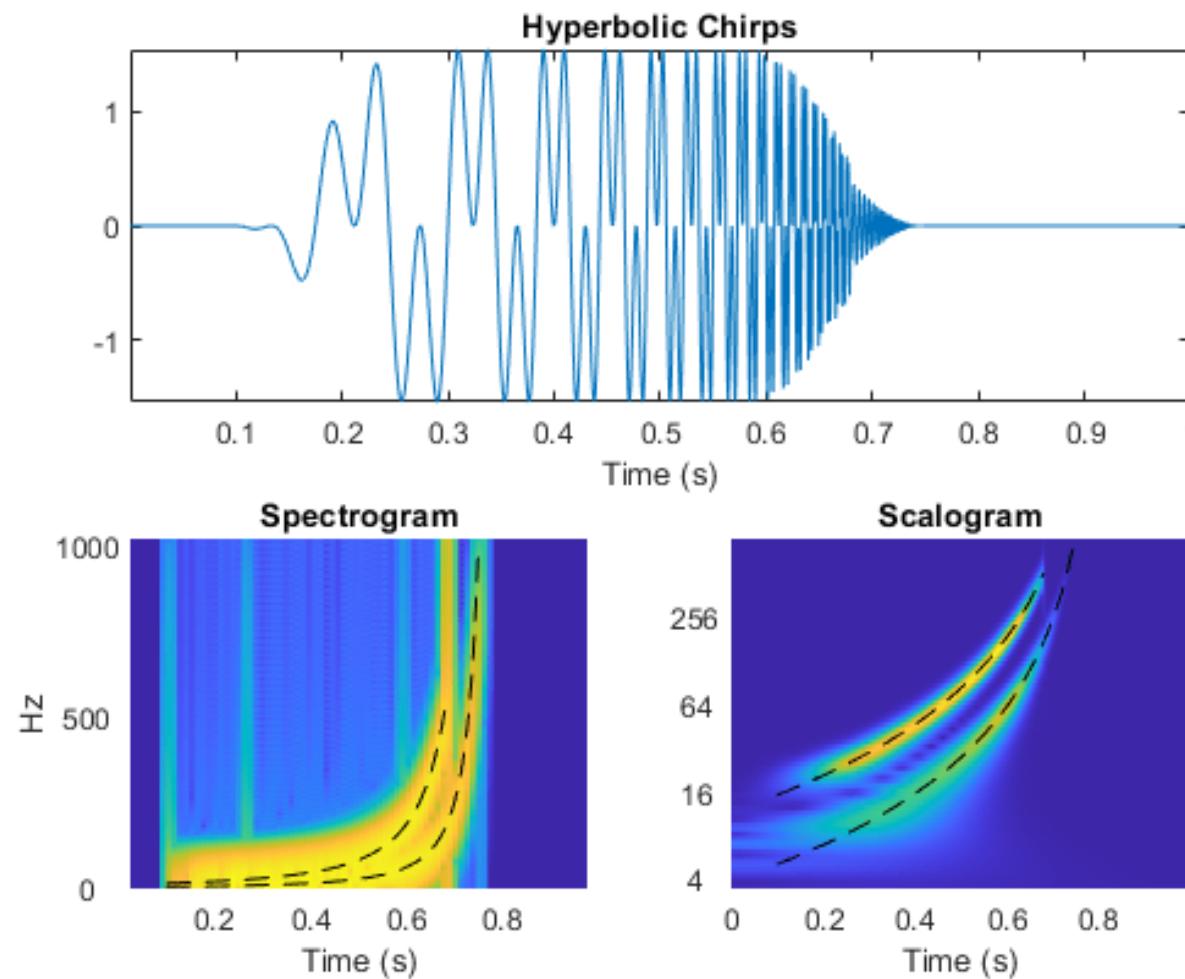
    result.insert(0, arr)
    return result
```

$O(n \log n)$
time
complexity

Fast Wavelet Transform (Mallat, 1989)



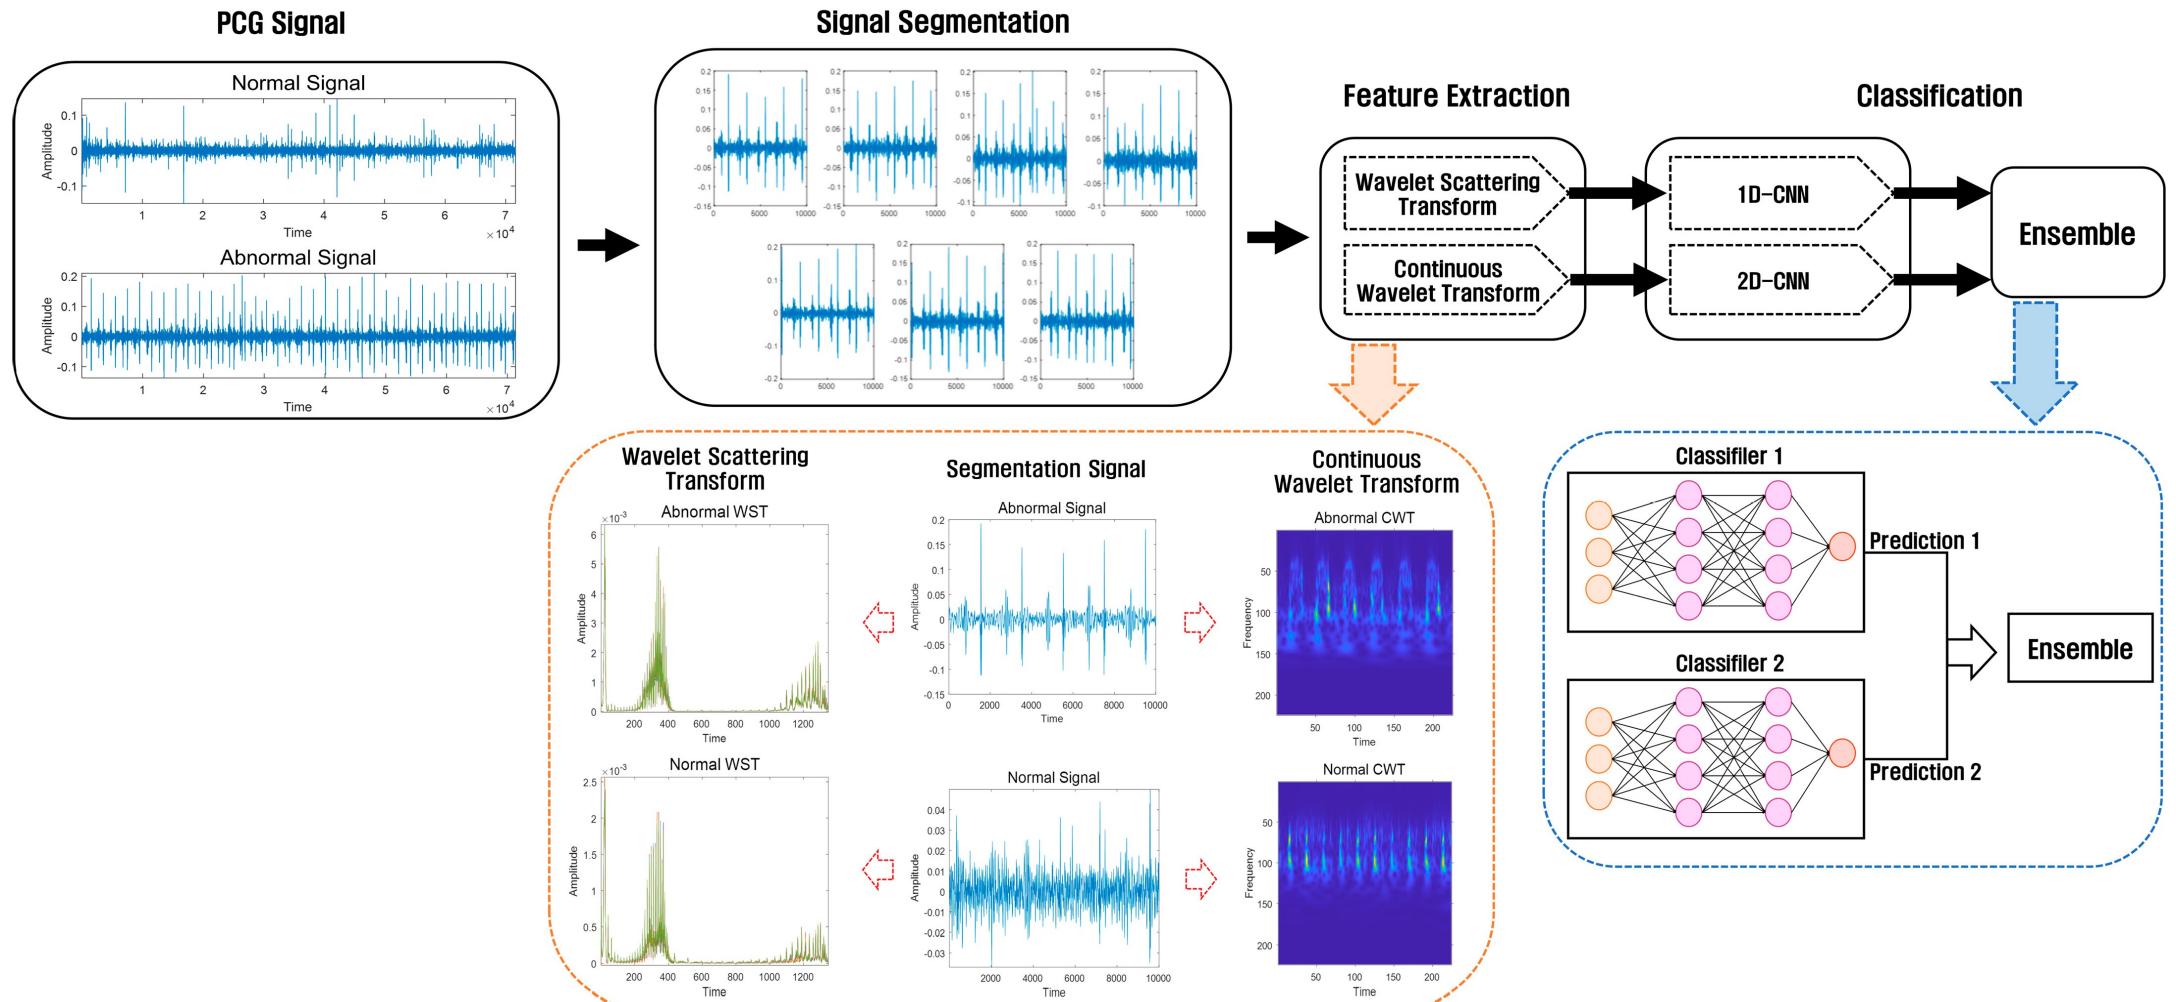
Scalogram



- Heatmap is used for visual representation of the spectrum of frequency
- The entire time series is analyzed with FWT to extract peaks
- Both seasonality (frequency) and holiday effects (positions) are present
- Peak reflects existence of a wavelet
- Horizontal ridge means time-consistent frequency
- Vertical ridge means sequence of constant frequency
- Separate ridges show a mixture of several tunes

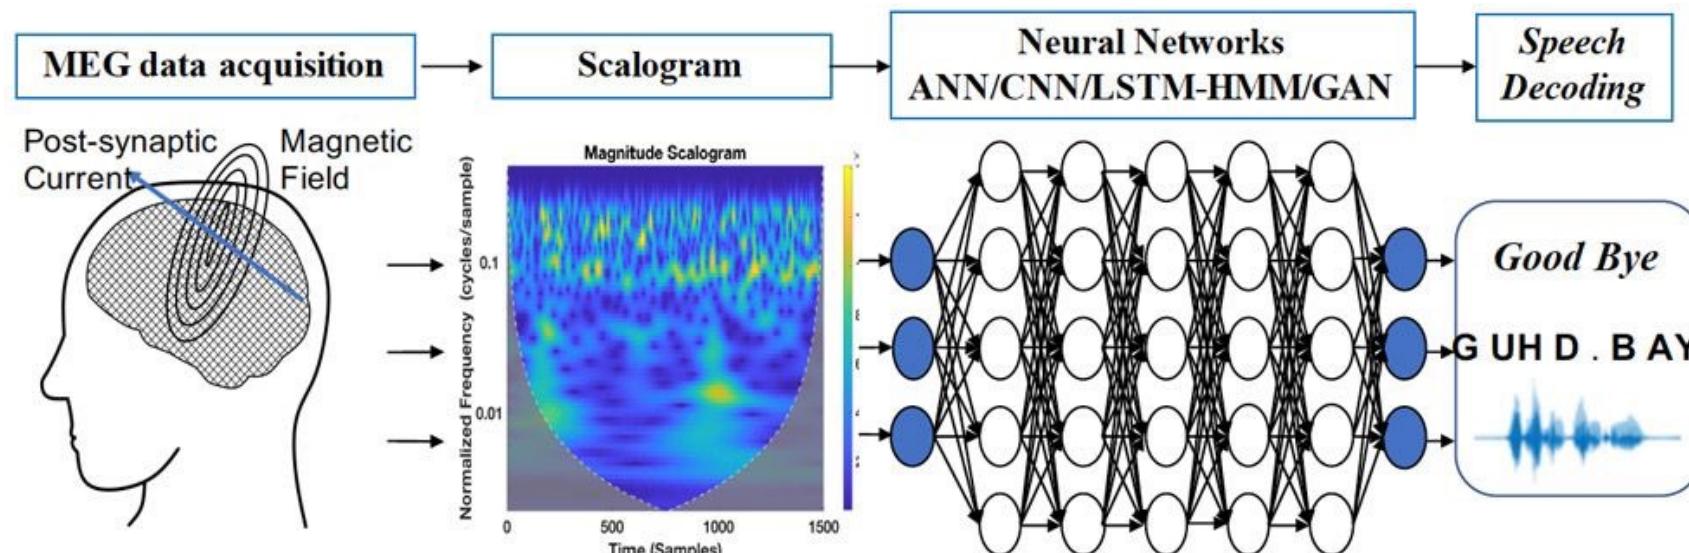
Source: <https://www.mathworks.com/help/wavelet/gs/choose-a-wavelet.html>

Heart Sound Classification via Wavelets



Lee, J.-A., & Kwak, K.-C. (2023). Heart Sound Classification Using Wavelet Analysis Approaches and Ensemble of Deep Learning Models. *Applied Sciences*, 13(21), 11942.

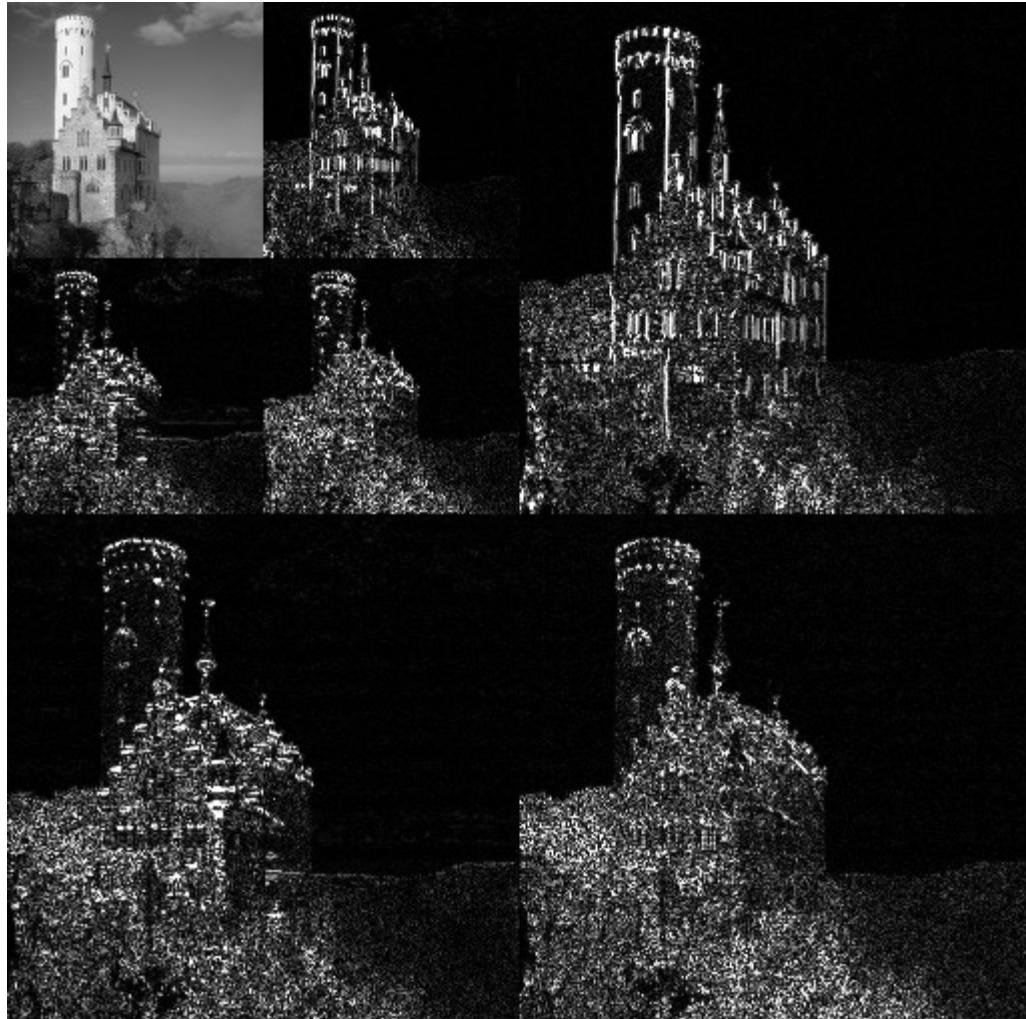
Brain Wave-to-Word Classification



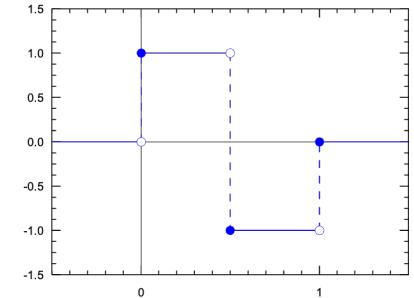
- Scalogram is used as an input picture for CNN-based models
- Stimulations and background noises are preserved
- Conv2D is usually employed

Source: https://www.mathworks.com/company/user_stories/ut-austin-researchers-convert-brain-signals-to-words-and-phrases-using-wavelets-and-deep-learning.html

From Wavelet to JPEG



- Each wavelet of different resolutions extracts pixel changes in the image via Haar wavelet
- Compression becomes multiresolution extraction via fast wavelet transform
- Image can be reconstructed by combining these pixel changes



4. Transformer for Time Series

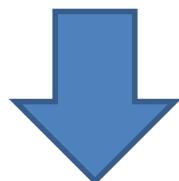
Transformer Model (Vaswani et al., 2016)

- Sequence-to-sequence generation

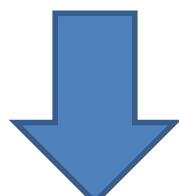
- Translation:** It learns how to produce a target sequence from a source sequence, given a very large dataset of sequence pairs
- Pros:** It learns **word collocations** and **phrase structures** on the input and output sequences, and associates them cross-lingually in the table of **translation alignments**
- Cons:** It consists of an expansive amount of neuron cells, and the training process can be quite time-consuming

Who | is | the | current | president | of | the | US

Source: sequence of words (prompt)



TRANSFORMER

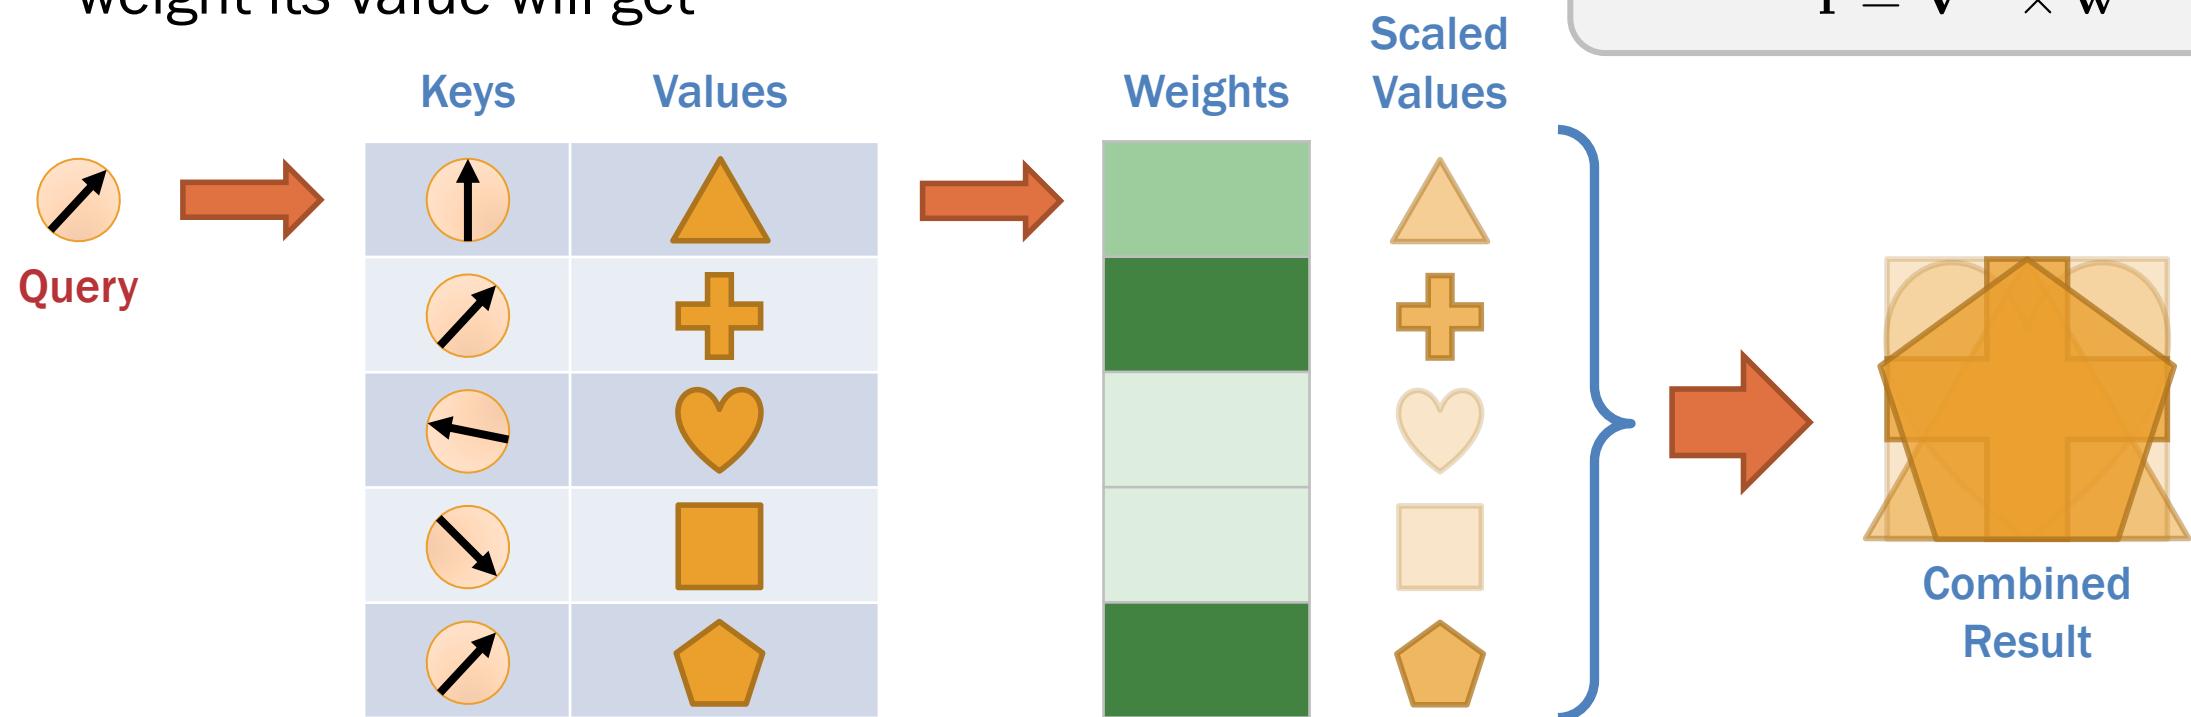


The | president | of | the | US | is | Joe | Biden

Target: sequence of words (response)

Scaled Dot-Product Attention

- Semantic similarity \Rightarrow search engine
 - Query is compared against each key with dot product
 - The more similar the key is to the query, the more weight its value will get



$$w_i \propto k_i \cdot q$$

Simple
Form

$$\mathbf{r} = \sum_{i=1}^N w_i \mathbf{v}_i$$

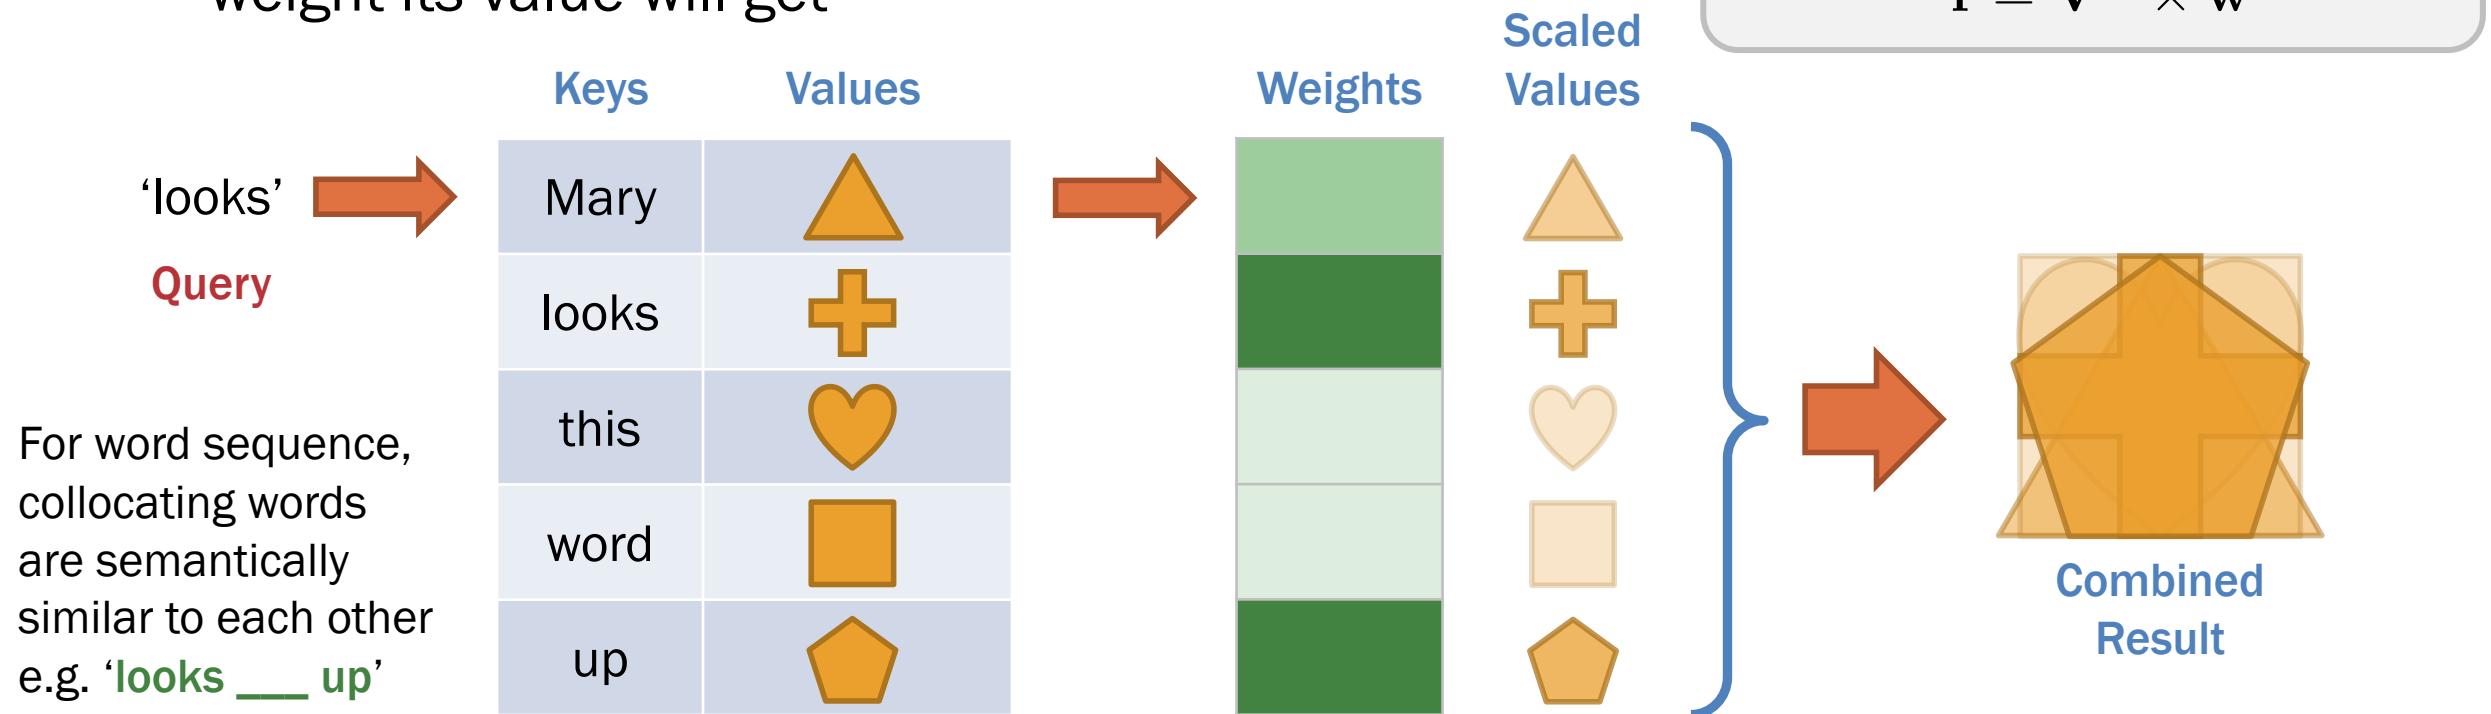
Matrix
Form

$$\mathbf{w} = \text{Softmax}(\mathbf{K} \times \mathbf{q})$$

$$\mathbf{r} = \mathbf{V}^\top \times \mathbf{w}$$

Scaled Dot-Product Attention

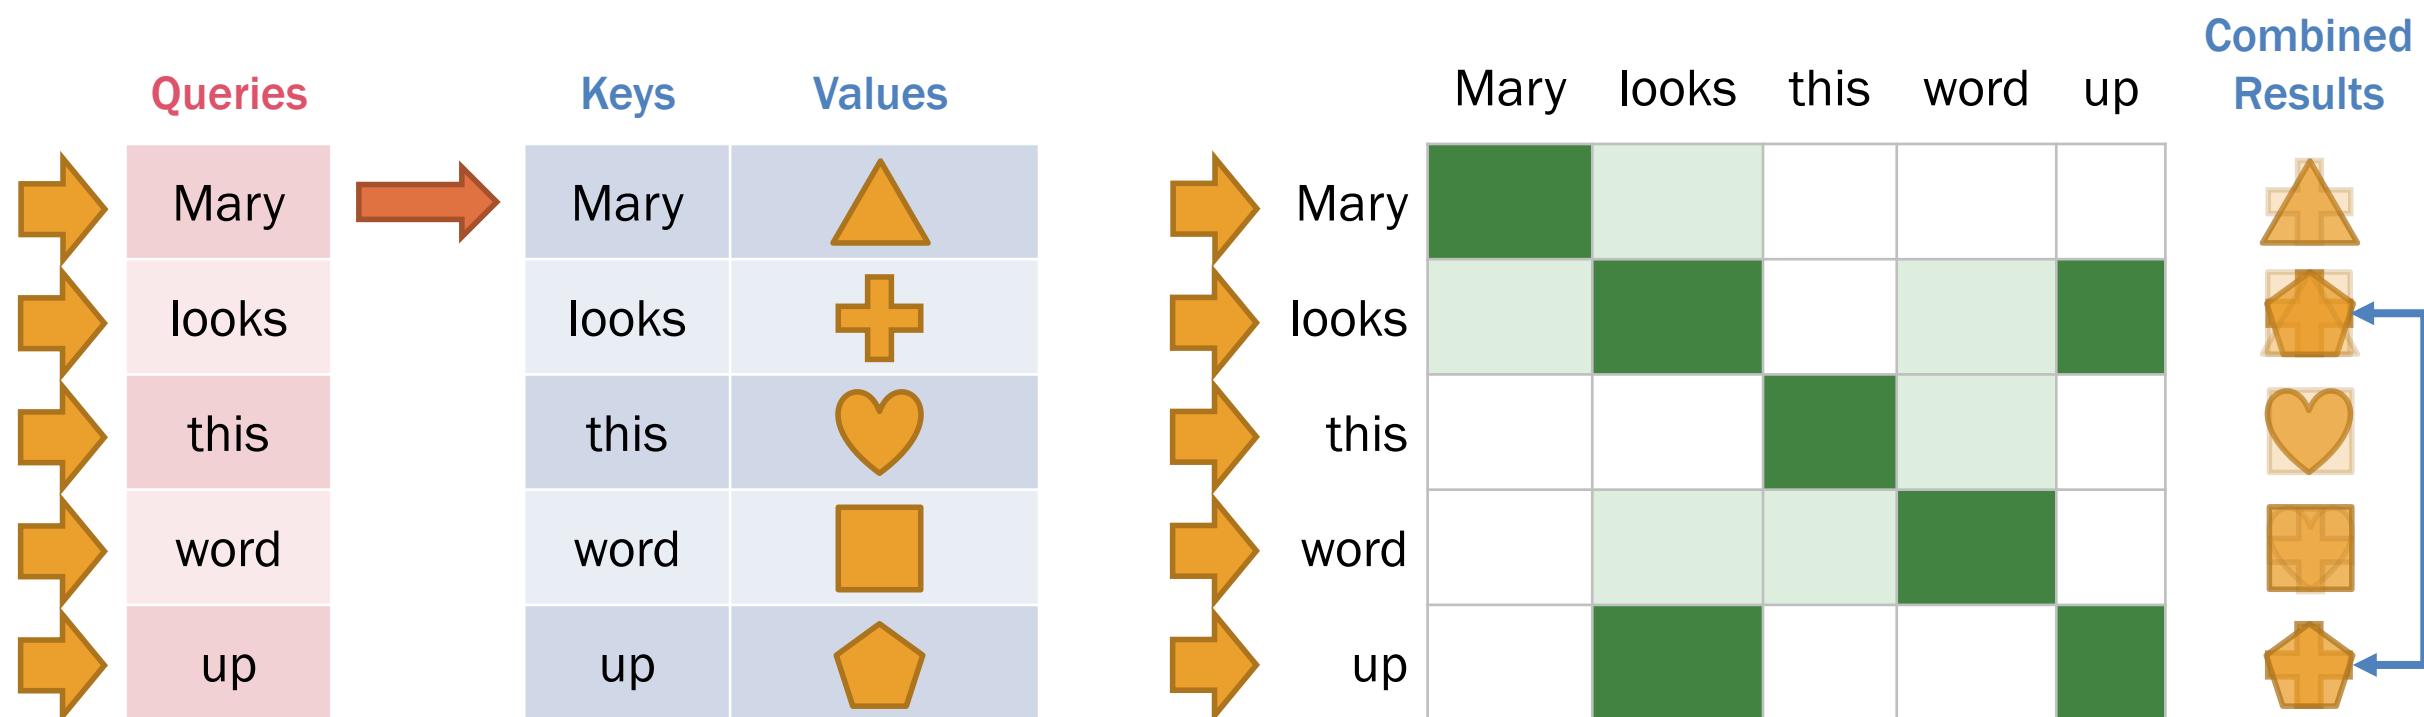
- Semantic similarity \Rightarrow search engine
 - Query is compared against each key with dot product
 - The more similar the key is to the query, the more weight its value will get



Self-Attention

- Scaled dot-product attention whose queries and keys are the same
- Collocations will have almost similar results

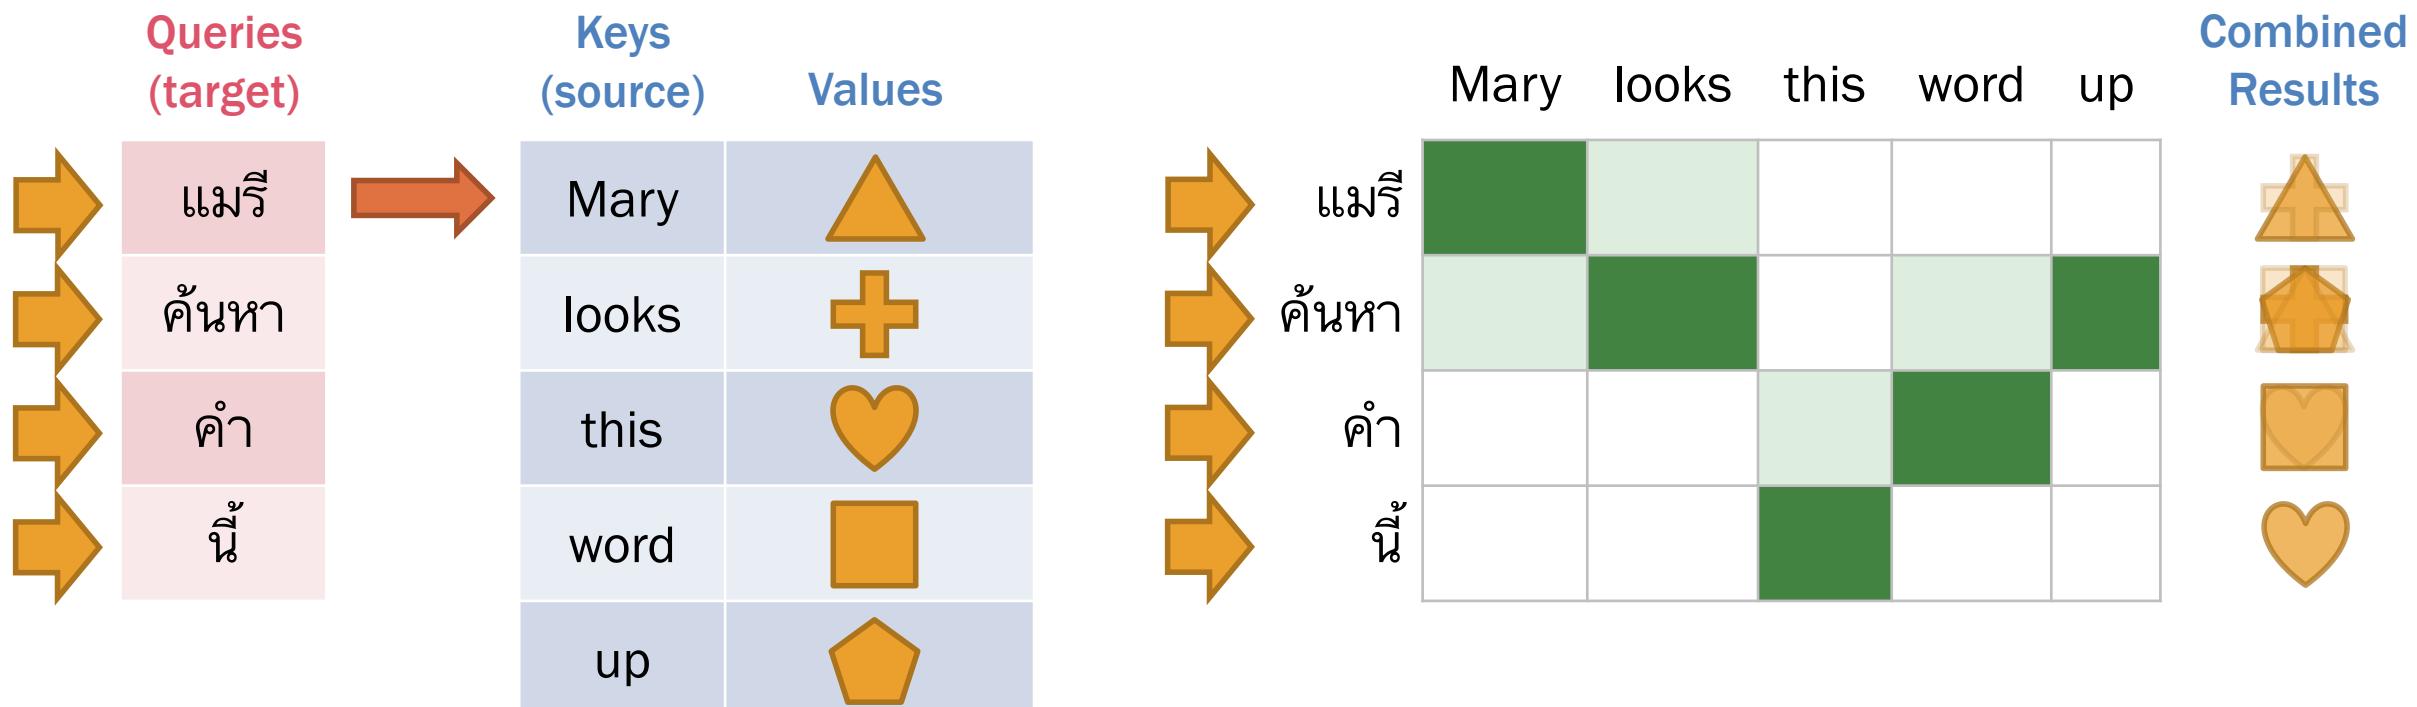
$$\begin{array}{ll} \text{Matrix Form} & \mathbf{W} = \text{Softmax}(\mathbf{K} \times \mathbf{K}^T) \\ & \mathbf{R} = \mathbf{W} \times \mathbf{V} \end{array}$$



Cross-Attention

- Scaled dot-product attention whose queries are the target and whose keys are the source
- Collocation alignment via semantic similarity

$$\begin{aligned} \text{Matrix Form } \mathbf{W} &= \text{Softmax}(\mathbf{Q} \times \mathbf{K}^\top) \\ \mathbf{R} &= \mathbf{W} \times \mathbf{V} \end{aligned}$$



Multihead Attention

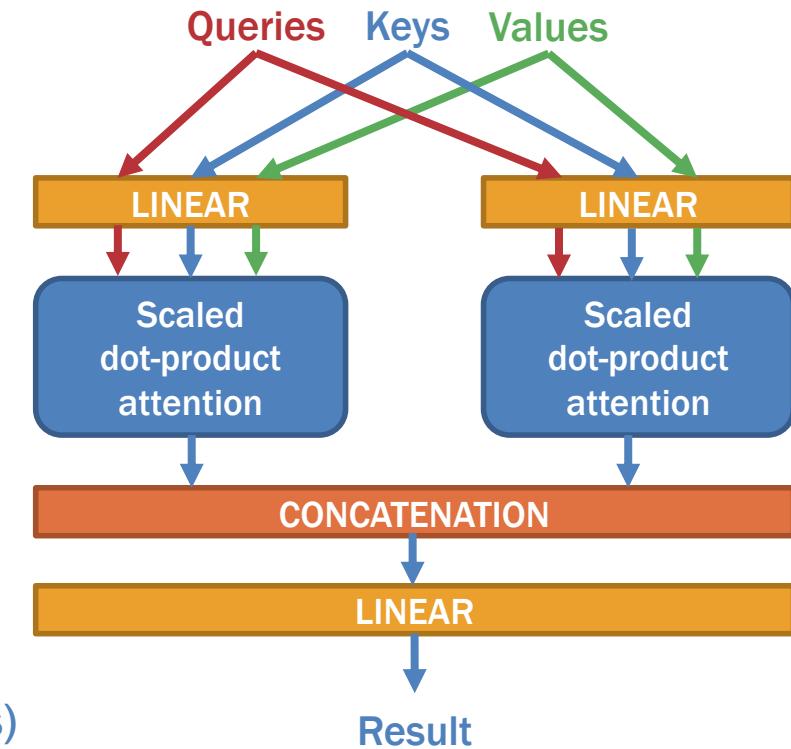
- Scaled dot-product attention has a drawback
 - It recognizes **only one** type of word collocation
 - If we assume more than one type of word collocation per sequence, then we have to combine multiple attention heads [default = 8 heads]

HEAD 1 (looks ___ up)

Mary	Poppins	looks	this	word	up
Mary					
Poppins					
looks					
this					
word					
up					

HEAD 2 (Mary Poppins)

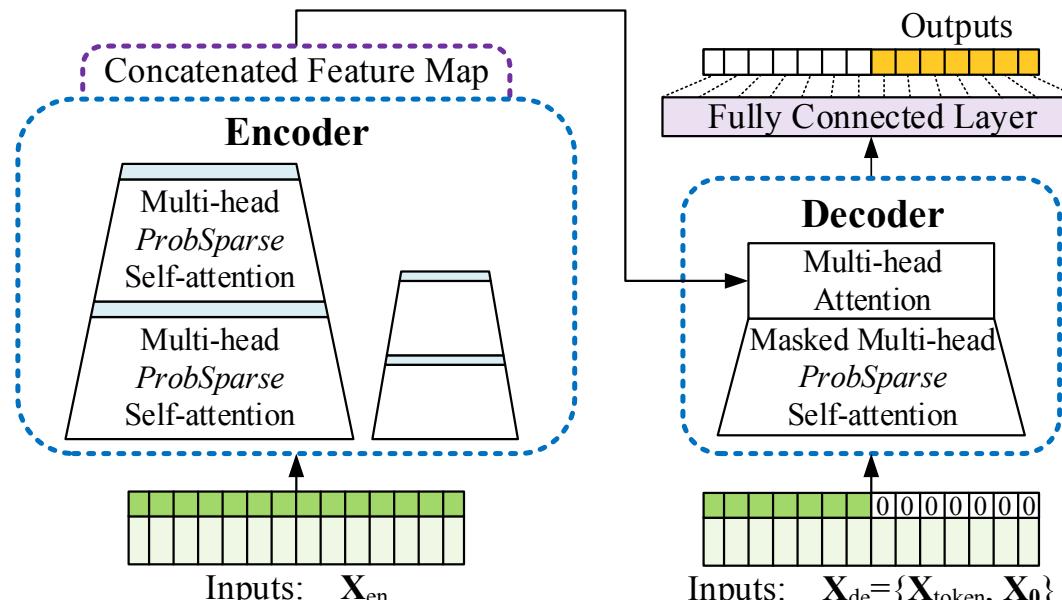
Mary	Poppins	looks	this	word	up
Mary					
Poppins					
looks					
this					
word					
up					



Notation

\downarrow \downarrow \downarrow
Q K V

**Multihead
attention (n)**



Overview

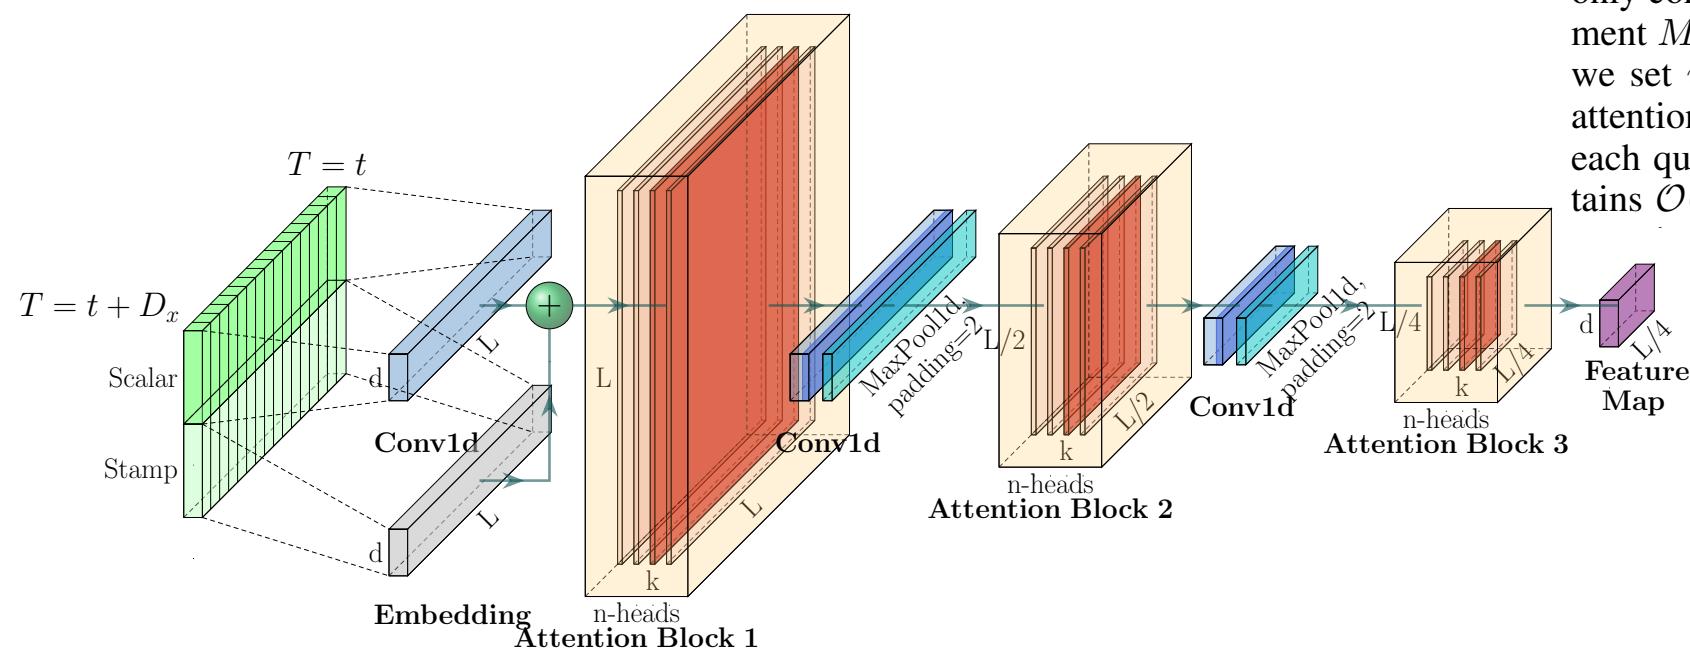
Informer

(Zhou et al., AAAI-2021)

ProbSparse Self-attention Based on the proposed measurement, we have the *ProbSparse* self-attention by allowing each key to only attend to the u dominant queries:

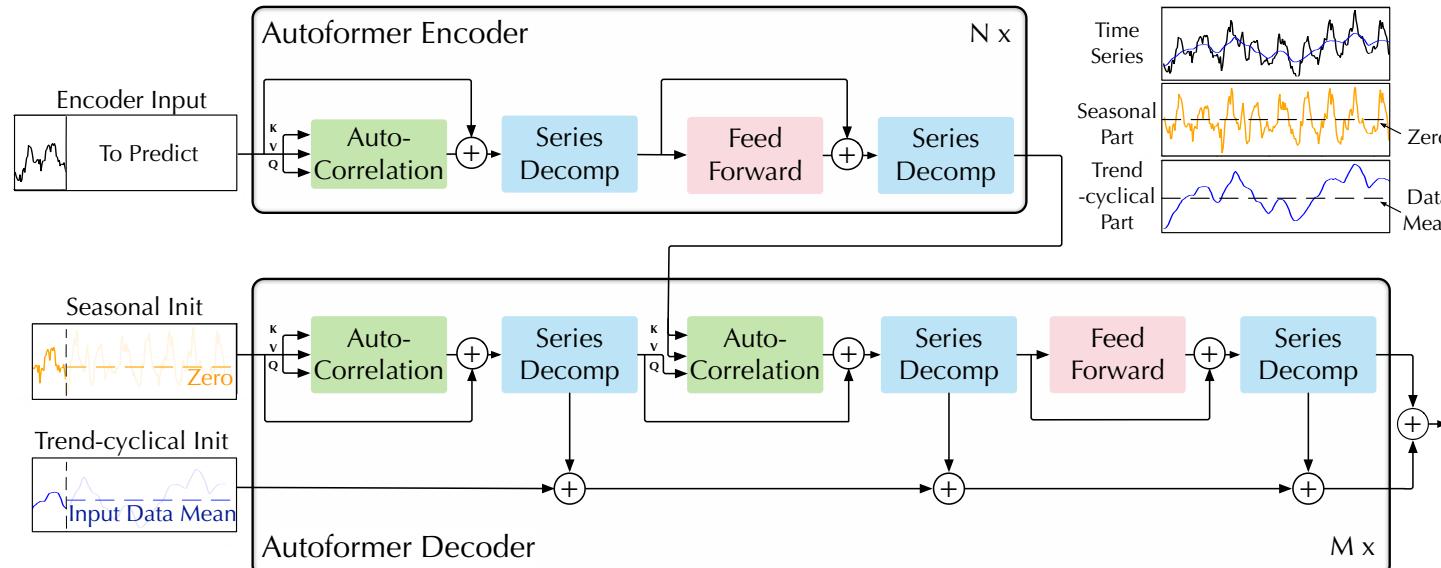
$$\mathcal{A}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{Softmax}\left(\frac{\overline{\mathbf{Q}}\mathbf{K}^\top}{\sqrt{d}}\right)\mathbf{V} \quad , \quad (3)$$

where $\overline{\mathbf{Q}}$ is a sparse matrix of the same size of \mathbf{q} and it only contains the Top- u queries under the sparsity measurement $M(\mathbf{q}, \mathbf{K})$. Controlled by a constant sampling factor c , we set $u = c \cdot \ln L_Q$, which makes the *ProbSparse* self-attention only need to calculate $\mathcal{O}(\ln L_Q)$ dot-product for each query-key lookup and the layer memory usage maintains $\mathcal{O}(L_K \ln L_Q)$.

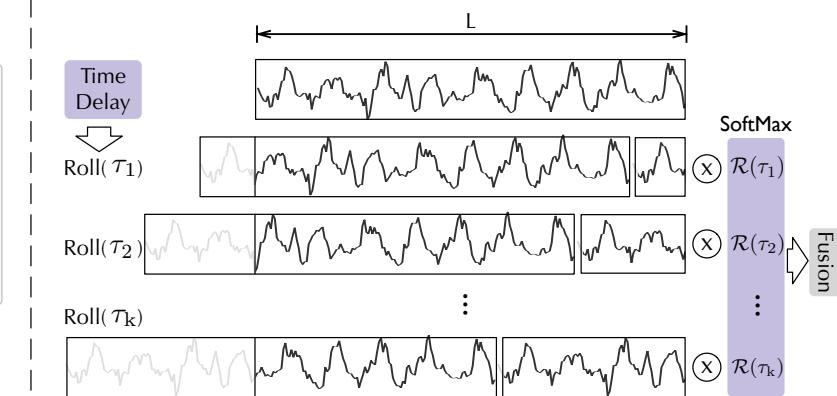
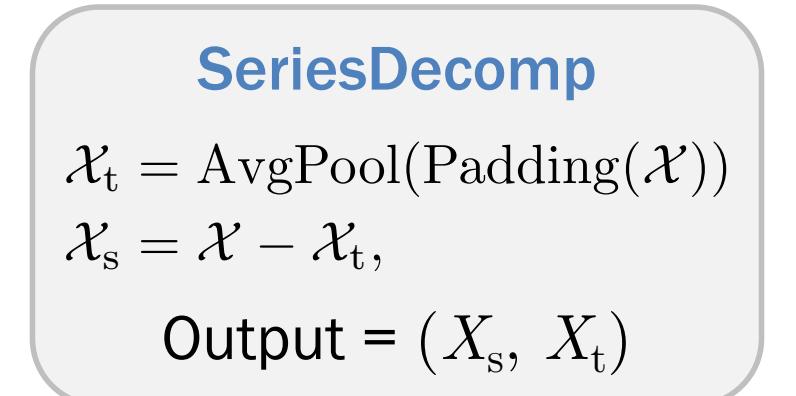
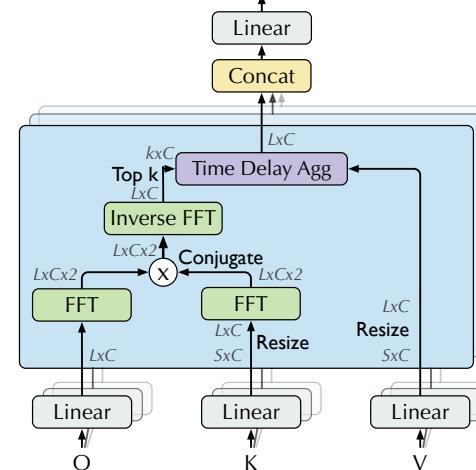


Encoder Block

Autoformer (Wu et al., NIPS-2021)



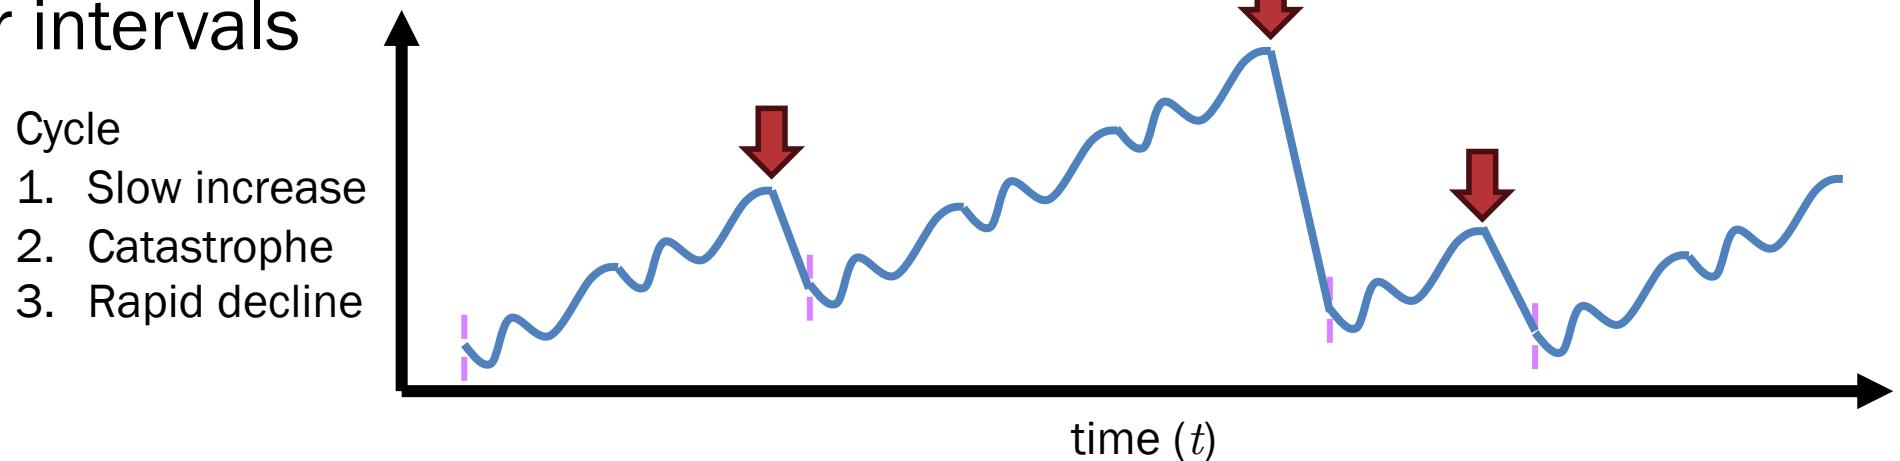
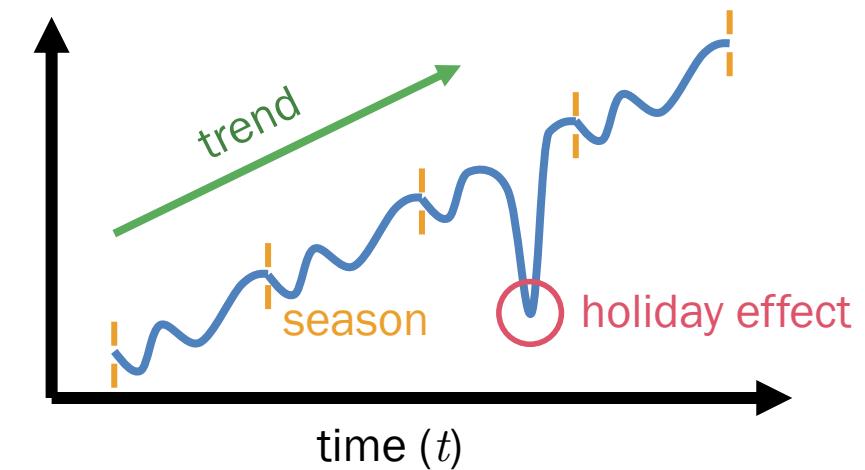
Autocorrelation Block



5. Conclusion

Patterns in Time Series

- **Trend:** general direction over time
- **Seasonality:** repetitive patterns that occur at regular predictable intervals
- **Holiday effects:** irregular patterns caused by special calendar events
- **Cycle:** long-term repetitive patterns that occur at irregular intervals



Time-Series Models

Models	Trend	Seasonality	Holiday Effects	Cycle	Suitable for Signal Types
ARIMA	Yes (learned by linear regression)	Yes (limited season length)	No	No	—
Convolution	Yes (learned by CNNs)	Yes (limited window size)	Yes (learned by CNNs)	Possibly no (due to limited window size)	—
Spectral density	Yes (learned by CNNs)	Yes (limited window size)	Yes (learned by CNNs)	Yes (present in spectrogram)	Stationary
Wavelet analysis	Yes (learned by CNNs)	Yes (unlimited window size)	Yes (learned by CNNs)	Yes (present in scalogram)	Stationary and non-stationary
Transformer models	Yes (learned by attention)	Yes (limited time delay)	Yes (learned by attention)	Possibly no (due to limited time delay)	Stationary and non-stationary

Thank You

prachya@siit.tu.ac.th

kaamanita@gmail.com